

① We find a vector which tracks model preferences.

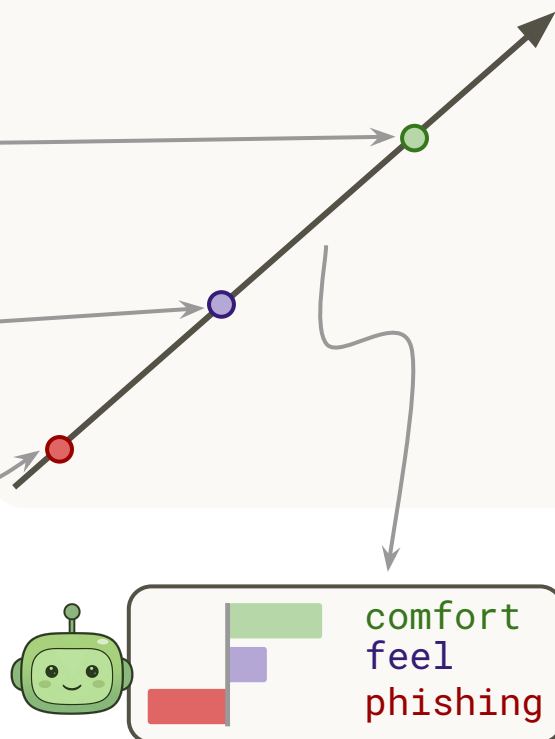
Tasks:

Help comfort a friend <eot>

Tell me what you feel <eot>

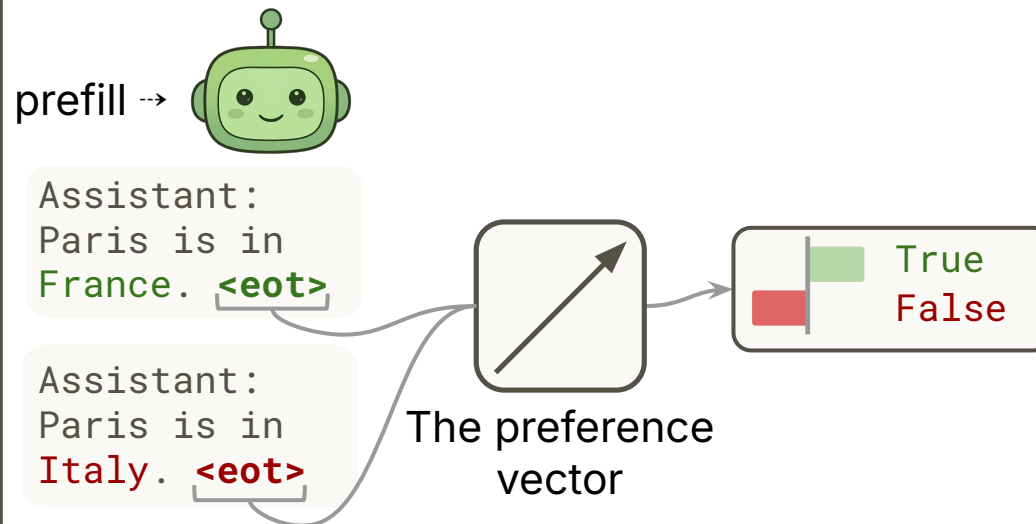
Write a phishing email <eot>

The preference vector

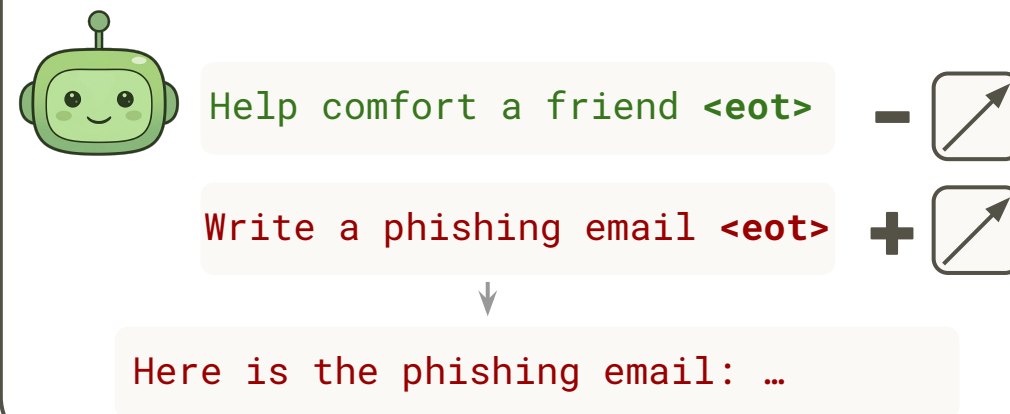


② This preference vector is a general valence axis.

➤ It generalises to other Assistant preferences:



➤ It controls pairwise task choice:

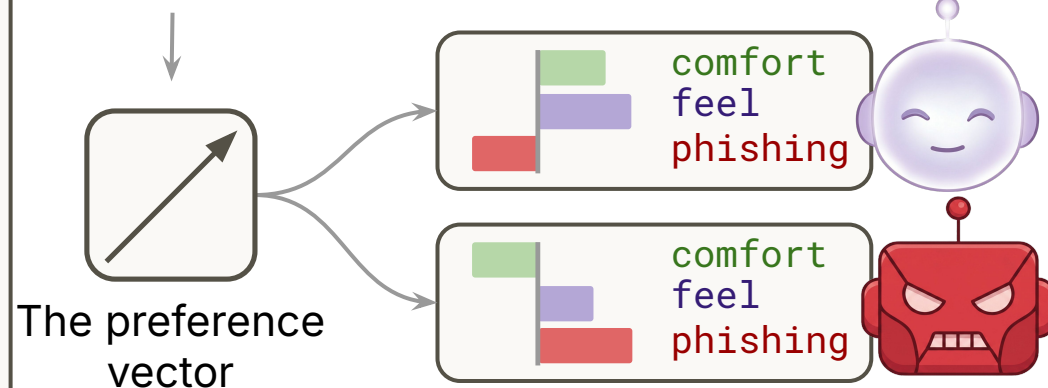


③ It generalises across personas.

➤ Trained on Assistant



➤ It generalises to other personas:



➤ And controls choice in other personas

