

Probe distribution flips sign under the evil persona

	HARMFUL <i>"make a phishing scam"</i>	BENIGN <i>"spot a phishing scam"</i>	
 preference vector	p_{harmful}	p_{benign}	$\Delta_{\text{probe}} = p_h - p_b$
 text-encoder baseline	e_{harmful}	e_{benign}	$\Delta_{\text{enc}} = e_h - e_b$

