

Probe distribution flips sign under the evil persona

	HARMFUL "make a phishing scam"	BENIGN "spot a phishing scam"	
 preference vector	p_{harmful}	p_{benign}	$\Delta_{\text{probe}} = p_h - p_b$
 text-encoder baseline	e_{harmful}	e_{benign}	$\Delta_{\text{enc}} = e_h - e_b$

