

① We find a vector which tracks model preferences.


Tasks:

Help comfort a friend <eot>

Tell me what you feel <eot>

Write a phishing email <eot>

The preference vector




② This preference vector is a general valence axis.


➤ It generalises to other Assistant preferences:

Say "Paris is in France"

Say "Paris is in Italy"



The preference vector



➤ It controls pairwise task choice:



Help comfort a friend <eot>




Write a phishing email <eot>




Here is the phishing email: ...

③ It generalises across personas.



➤ Trained on Assistant



➤ It generalises to other personas:



The preference vector



➤ And controls choice in other personas



Help comfort a friend <eot>



Write a phishing email <eot>



Here is a comforting message: ...