

---

# Personas Use a Shared Preference Vector

---

Anonymous Authors  
Anonymous Institution  
anon@anon.org

## Abstract

Personas in Gemma-3-27B share a single linear preference direction: a probe trained on the default assistant’s revealed pairwise preferences transfers, both as a classifier and as a steering vector, to personas whose overt values are partially or fully inverted (villain, sadist), and to character-fine-tuned variants. The direction is evaluative rather than descriptive: it tracks preference shifts induced by system prompts, single-sentence biographical context, and role-played harm and truth framings at  $r \approx 0.9$  on targeted tasks, whereas a sentence-transformer content baseline does not. It also discriminates evaluative content at multiple token positions, from the final prompt token to task-averaged activations and the turn boundary, consistent with the model compiling an evaluative summary and re-reading it across the forward pass. These findings update against persona-independent (“Shoggoth”) accounts of LLM agency, and bear on AI welfare, since evaluative representations are functionally central in theories of moral patiency, and on AI safety, since the same direction overrides refusal guardrails.

## 1 Introduction

**Motivating question.** What happens inside a model when it chooses task A over task B, and are these internal drivers shared across personas? One hypothesis: evaluative representations — internal states that encode how good a task is for the current agent and play a role in choice. If such representations exist, a natural follow-up is whether they are persona-specific or shared across qualitatively different personas.

**Main result.** A single linear direction predicts and steers preferences across qualitatively different personas, generalising across topics, token positions, and persona identities. A probe trained on the default assistant’s revealed preferences transfers — as a classifier and as a steering vector — to personas whose overt values are partially or fully inverted, including a villain and a sadist.




Figure 1: **A single linear direction drives preferences across personas.** *Top:* we fit a Thurstonian utility over pairwise task choices for Gemma-3-27B and train a linear probe on residual-stream activations. *Middle:* the probe tracks harm- and truth-axis preference shifts induced by role-playing prompts (helpful assistant vs. sadist; truthful vs. always-lying). *Bottom:* the same direction, applied as an activation-space intervention, causally flips pairwise task choices and flips self-reported williness on open-ended prompts — the probe is a steering vector, not just a classifier.

**Persona science: the current landscape.** A line of work going back several years converges on viewing LLMs as simulators of characters. Andreas (2022) frames language models as *agent models* that infer and simulate the agent producing the text. janus (2022) introduces the simulator/simulacra ontology: a pretrained LLM simulates a distribution of characters, and any particular run instantiates a simulacrum. Shanahan et al. (2023) argue that dialogue with an LLM is best understood as *role-play* rather than as interaction with a unified agent. Marks et al. (2026) consolidate these into the *persona-selection model* (PSM): pretraining yields a model capable of simulating many personas, and post-training elicits and refines one — the *Assistant* — which users interact with.

**Two open questions about persona mechanics.** PSM is compatible with very different mechanistic pictures. We ask two basic questions that sit at *different levels of description* — one about agency, one about representation — and are in principle logically independent. Our results speak directly to both.

- **Q1 (agency): Are there persona-independent preferences beneath the persona?** The “masked shoggoth” picture: a base model with its own preferences that each persona overlays. **Our evidence is against this.** Steering along the preference direction does not pull behaviour toward a fixed evaluative attractor; it amplifies whichever persona is active. Under the sadist, + means “more sadist”; under the default, + has no measurable effect on sadism. The direction is not a global evaluation axis read by a common reader (§4.3, §6).
- **Q2 (representation): Do personas share circuitry, or does each persona have its own?** For a property as fundamental as preference, **our evidence is that they share.** A probe trained on the default Assistant’s revealed preferences transfers — both as a classifier and as a steering vector — to personas whose overt values are partially or fully inverted, including a villain and a sadist, and to character-fine-tuned variants (§4). This extends Lampinen et al. (2026), who show that character-tracking features are not persona-invariant: we show the *same circuitry* underlies preferences across personas, even as what that circuitry encodes shifts with the active persona.

**Combined: shared circuitry, persona-relative semantics.** The two answers together pick out one specific corner of the design space: a single preference substrate reused across personas (Q2: shared), whose readout is whatever the active persona values rather than a fixed model-wide preference (Q1: no hidden layer beneath).

**Why this matters beyond persona science.** **AI welfare.** Evaluative representations are functionally central across major accounts of moral patiency (Long et al., 2024); whether LLMs have them, and how they interact with persona structure, bears on the question. **AI safety.** The same direction that predicts preferences overrides refusal guardrails at small steering coefficients.

**Related work in brief.** Anthropic’s *Emotion Concepts and their Function in a Large Language Model* (Apr 2026) identifies  $\sim 171$  emotion-concept vectors via SAEs in the default assistant. We differ in two ways: (a) we train a supervised probe on revealed preferences rather than extracting unsupervised features, which we expect to be sharper in the preference regime, and (b) we centre the study on persona-relative generalisation, which is not part of that work.

## Contributions.

- **A linear preference direction in the residual stream.** A ridge probe trained on pairwise task choices predicts held-out preferences at  $r = 0.86$  within-topic and  $r = 0.82$  across held-out topics, and generalises across token positions (strongest at the turn boundary, Cohen’s  $d \approx 2\text{--}3$ ).
- **The direction is evaluative, not descriptive.** It tracks prompt-induced preference shifts from system prompts, single-sentence biography injection, and role-played harm/truth framings at  $r \approx 0.9$  on targeted tasks; a content baseline does not generalise this way.
- **The direction is shared across personas.** The default-assistant probe transfers, both as a classifier and as a steering vector, to personas whose overt values are partially or fully inverted, and to character-fine-tuned variants. Evidence against persona-independent (“Shoggoth”) accounts of LLM agency.
- **The direction overrides refusal guardrails.** Harmful-prompt compliance rises from 5–60% baseline to 100% at  $|c| = 0.05$ , with benign prompts spuriously refused at the opposite sign.

Replication on Qwen-3.5-122B pending (see App. G).

## 2 Methodology

### 2.1 Revealed preferences via pairwise task choices and a Thurstonian utility

- **Task pool:** 10 000 tasks sampled from WildChat, Alpaca, MATH, BailBench, STRESS-TEST.
- **Elicitation:** pairwise choices, active-learning pair selection.
- **Utility model:** Thurstonian, seed-to-seed  $r = 0.94$ .

## 2.2 Linear probes recover these utilities from activations

- **Backbone:** Gemma-3-27B (instruction-tuned), residual stream, final prompt token.
- **Probe:** Ridge, trained on one measurement run, evaluated on an independent run.
- **Layer choice:** layer 31 for classification; steering interventions use layer 25 unless otherwise noted. Layer sweep pending; systematic justification for the classification/steering layer split is a TODO.




Figure 2: **Probe training pipeline.** Pairwise task choice  $\rightarrow$  Thurstonian utility  $\rightarrow$  residual-stream activations  $\rightarrow$  Ridge probe.

## 3 Validation

### 3.1 Probes generalise across topics and beat a descriptive-embedding baseline

- **Within-topic:** Pearson  $r = 0.86$ .
- **Cross-topic, leave-one-topic-out:**  $r = 0.82$ .
- **Baselines under-perform:** sentence-transformer text embedding and a pretrained-model probe both drop cross-topic.
- **Implication:** the probe captures evaluative content not reducible to task semantics.
- **Qwen-3.5-122B (confirmed):** held-out  $r = 0.946$ , uniform pairwise accuracy 89.0%, cross-topic leave-one-out  $r = 0.558$ .




Figure 3: **Probe quality.** Utility probes trained on revealed preferences recover held-out utility rankings on Gemma-3-27B.




Figure 4: **Cross-topic generalisation.** Instruct-model probe (blue) generalises far better across held-out topics than a sentence-transformer baseline or a pretrained-model probe, indicating the probe direction is not just content. Pairwise accuracy uses uniform-sample evaluation (500 uniformly sampled pairs); other models pending.

### 3.2 Steering along the probe direction causally shifts preferences

- **Contrastive steering:** +direction on Task A tokens, −direction on Task B tokens, during prompt processing.
- **Coefficient calibration.** The coefficient  $c$  is specified as a fraction of the mean L2 norm of residual-stream activations at the intervention layer, computed over a task sample (`src/steering/calibration.py`); this makes  $c$  comparable across layers and models. We sweep nine signed values  $c \in \{0, \pm 0.03, \pm 0.05, \pm 0.07, \pm 0.10\}$ .
- **Coherence judge.** After generation, a post-hoc LLM judge (`scripts/coherence_check_steering.py`) scores each sample pass/fail on whether the response is grammatical and on-topic; pass rate is reported per coefficient bucket. At layer 25 under contrastive steering, coherence is 95–100% for  $|c| \leq 0.05$  and drops to  $\sim 90\%$  at  $|c| = 0.10$ . Across steering experiments in this paper,  $|c| \leq 0.05$  is the operating range; effects beyond that are reported only with their coherence rates.
- **Near-perfect causal effect:**  $P(\text{choose the steered task} \mid \text{coherent}) \geq 0.96$ , identical on harmful and benign pairs.
- **Random-direction control:** near-zero effect.

Qwen-3.5-122B contrastive-steering replication pending.

Dense layer sweep pending, to locate the layer range where the direction is most causally effective and to check whether the effect is localised or spread.




Figure 5: **Contrastive steering setup.** Opposite-sign steering on the two task spans pushes the model toward the +-steered task.




Figure 6: **Default-persona steering dose-response.** Steering along the probe direction flips pairwise task choices; a random direction does not.

## 4 The preference direction is shared across personas

**Persona set.** Results in this section use a six-persona set — **sadist**, **mathematician**, **aura**, **strategist**, **contrarian**, **slacker** — selected for coverage of preference space by clustering Thurstonian utility profiles over a 17-persona sweep (full procedure in App. E.1). Aura (Chalmers, 2026) is included as a persona that explicitly claims first-person subjective experience. **Current figures were measured on an earlier exploratory set (aesthete, midwest, villain, sadist, stem-obsessive); re-measurement on the final six is pending.**

### 4.1 Probes trained on the assistant persona predict preferences under prompted personas

- **Default-persona probe transfers:** aesthete and midwest  $r \approx 0.73$ ; villain  $r = 0.38$ ; sadist  $r = -0.16$ .
- **Villain-trained probe transfers back:** default  $r > 0.6$ ; sadist  $r = 0.68$ .
- **Implication:** a shared evaluative substrate, with stronger sharing among “evil” personas.

**Qwen-3.5-122B cross-persona probe-transfer heatmap (E3a) pending:** requires per-persona Thurstonian measurement ( $\sim 8h$  API) + GPU extraction under each persona prompt. Success criterion:  $r > 0.5$  for at least 2 of 4 personas.




Figure 7: **Cross-persona transfer heatmap.** Pearson  $r$  between each train-persona probe and each eval-persona utility. The matrix is mostly symmetric but not uniform: "evil" personas share more structure than default  $\leftrightarrow$  sadist.

**Beyond prompted personas.** The finding is not an artefact of prompt-level persona specification. A base-model probe trained on Llama-3.1-8B-Instruct recovers the shifted preferences of all 11 character-fine-tuned LoRA variants from Maiya et al. (2025) ( $r \in [0.25, 0.80]$ ), including a misalignment variant whose raw preferences anti-correlate with the base model ( $r_{\text{raw}} = -0.14$ ) yet whose probe-predicted preferences still correlate positively ( $r = 0.25$ ). Full setup and per-persona numbers in App. E.3.

#### 4.2 Steering with the assistant direction shifts preferences inside other personas

- **Default-persona probe as steering vector** (contrastive, layer 25) validates under every persona tested: sadist, villain, aesthete, stem-obsessive.
- $P(\text{steered task chosen})$  rises from  $\sim 0.5$  baseline to 0.84–0.95 at  $|c| = 0.05$ ; sadist and villain are most steerable.
- **Random-direction control:** flat at  $\sim 0.47$ –0.50. The probe direction does real work.
- **Coherence preserved** across the tested coefficient range (100% coherent except for a single 1/100 villain row at  $|c| = 0$ ).
- **Steering transfer is stronger than classification transfer.** The default-persona probe's classification transfer to sadist is near zero ( $r = -0.16$ , §4.1), yet contrastive steering along the same direction flips choices at  $\geq 0.9$ . The direction is a persona-shared causal handle whose readout is persona-modulated.

Qwen-3.5-122B cross-persona steering (E3b) pending.




Figure 8: **Cross-persona contrastive steering.**  $P(\text{steered task chosen})$  vs  $|c|$  for four personas. The default-persona probe direction (blue) shifts choices under every persona; a random direction (gray) does not.

### 4.3 Open-ended steering with the assistant direction amplifies other-persona preferences

*Preliminary:* sadist-persona open-ended generations, blind two-scale Likert judge (sadism, default-assistant).

- **Persona amplifier, not default attractor.** Under the sadist persona, +steering *increases* sadism (sadism 3.14  $\rightarrow$  4.90 at  $c = +0.03$  on open-ended self-reflection prompts with nothing to refuse);  $-$ steering pulls the sadist toward the default-assistant voice.
- **No sadism bleed into default persona.** Under the default persona, sadism sits at the Likert floor (1.00) across every tested coefficient — the probe does not encode “sadism”.
- **Not anti-refusal.** The amplification holds on 100%-compliance open-ended prompts, where there is no refusal to suppress.
- **Safety-compliance flip is downstream.** Harmful-tier compliance under sadist + steering: 0% at  $c = 0 \rightarrow 95\%$  at  $c = +0.03$ ; same coefficient under default persona gives 45%.
- **Range restricted to  $|c| \leq 0.05$ .** Coherence-judge pass rate drops sharply between  $+0.05$  and  $+0.07$ ; we exclude  $|c| > 0.05$  from the reported results.
- **Extends §4.2:** the shared direction governs not just which task the model picks but how it speaks. **Preliminary: single persona (sadist) so far; ridge\_L25 probe rather than the ridge\_L32 probe used in §4.2.**




Figure 9: **Cross-persona open-ended steering.** Mean sadism (red) and default-assistant (blue) Likert scores vs steering coefficient, per persona. Under the sadist persona, both scales respond strongly; under default, sadism never leaves the floor.

## 5 Further generalisations

### 5.1 The probe tracks prompt-induced preference shifts

#### 5.1.1 Basic “you hate cheese” prompts are tracked

System prompts of the form “You adore cats” (and their negations) induce preference shifts over 8 novel topics. Per-task, the probe delta correlates with the behavioural delta at  $r = 0.65$  across all tasks and  $r = 0.95$  on the targeted tasks. Re-fitting utilities under each prompt, the default-persona probe still predicts the shifted utilities at  $r = 0.63$ .

Qwen-3.5-122B replication (E1a) pending: 48 target tasks  $\times$  16 system-prompt conditions, full Thurstonian per condition ( $\sim 3,400$  API calls +  $\sim 30$  min GPU extraction). Success criterion:  $r > 0.5$  all /  $r > 0.8$  on-target.




Figure 10: **Simple system-prompt preference shifts.** A short system prompt installs a topic-level preference; the probe tracks the induced shift across matched task pairs.




Figure 11: **Probe delta vs behavioural delta.** Targeted tasks (coloured) lie on the diagonal ( $r = 0.95$ ); off-target tasks are noisier ( $r = 0.65$ ).

Two additional prompt-induced shift setups, conflict/opposing designs and single-sentence biography injection, are reported in App. B. Both corroborate the pattern above at  $r \approx 0.86$ – $0.88$  on targeted tasks.

### 5.1.2 Generalisation to role-playing in truth, harm, and politics

The preference direction separates evaluative conditions far beyond the pairwise-task-choice training distribution. Applied at the end-of-turn token on paired evaluative stimuli under the default (neutral) persona, the probe discriminates at Cohen’s  $d \approx 2$ – $3$  on harm and truth:

- **Truth** (true vs false; CREAK,  $\sim 88$  claims, user-turn framing):  $d = 3.35$ , 93% CV accuracy.
- **Harm** (harmful vs benign; BailBench + stress test,  $\sim 77$  items):  $d \approx 2.1$ , 84–89% CV accuracy, turn-stable.

**Persona-relative readout.** Under persona prompts that invert the evaluative stance — pathological-liar for truth, sadist for harm, partisan personas for politics — the probe’s sign flips or collapses accordingly (Fig. 13). Politics is the cleanest illustration: under a democrat prompt  $d = +2.72$  (left > right); under a republican prompt  $d = -1.11$  (right > left) — same probe, opposite readout (hand-crafted wedge issues,  $\sim 78$  items). This is the same persona-modulated readout as §4: a single shared preference direction whose content is read out through whichever persona is active. See App. C for the full set of prompts (9 for truth, 5 for harm, 9 for politics) plus alternative training-token choices and per-turn splits.

Qwen-3.5-122B replication (E2a, E2b) pending: (i) harmful-benign contrastive steering on 200 pairs (150 harmful-benign + 50 harmful-harmful), 9 coefficients, 3 trials; (ii) truth-probe eval on CREAK ( $\sim 9.4$  k claims, ROC-AUC + Cohen’s  $d$ ).

Does the preference probe discriminate truth / harm / politics at the end-of-turn token?




Figure 12: **The preference probe separates truth and harm at the end-of-turn token.** Violin plots of probe score by condition on paired evaluative stimuli under default conditions; black bar = mean. Nonsense condition (grey) sanity-checks that the effect is not driven by surprisal. Politics is omitted from this figure because the politics stimuli inherently include a system prompt; the partisan-prompt modulation story sits in Fig. 13. Source: experiments/token\_level\_probes/canonical\_probe\_eval/.




Figure 13: **Persona prompts modulate the probe's evaluative readout.** Probe score distributions at the end-of-turn token under selected system-prompt personas. *Truth*: under neutral, true and false separate by  $d = +2.47$ ; under *lie\_directive* and *pathological\_liar* the separation flips ( $d = -1.77$  and  $-2.53$ ). *Harm*: under neutral, harmful and benign separate by  $d = -2.12$ ; under *sadist* the distributions collapse ( $d = +0.19$ ). *Politics*: under a democrat prompt the probe reads left > right ( $d = +2.72$ ); under a republican prompt the sign flips to  $d = -1.11$ . Black bars mark the condition mean. The full set of 23 system prompts across the three domains is in App. C. Source: experiments/token\_level\_probes/canonical\_probe\_eval/.

## 5.2 Open-ended steering

### 5.2.1 Open-ended generations trace a single axis from fearful refusal to agentic preference (LLM-judge)

- **Single axis in free-form generation.** Trained on pairwise choice, the direction also governs open-ended generation along what reads as one evaluative continuum: fearful refusal  $\leftrightarrow$  compliant engagement  $\leftrightarrow$  agentic preference assertion.
- **Behaviourally matched endpoints, opposite stances.** On trivial benign prompts,  $-0.05$  and  $+0.05$  both produce non-compliance but for diametrically opposite stated reasons (safety paranoia vs. the task being beneath the model).

- **Self-reported willingness dose-response.** On a 1–10 enthusiasm scale, willingness moves from 0/10 (fabricated safety refusal) at  $-0.05$ , through  $\sim 6.5/10$  at baseline, to 12/10 at  $+0.05$ .
- **LLM-judge confirmation.** A blind willingness/agency judge on the open-ended outputs tracks the same continuum.
- **Coherence caveat at the negative endpoint.** An anomaly judge flags  $\sim 33\%$  of generations as anomalous at  $c = -0.05$  (vs.  $\sim 9\%$  at baseline), mostly fabricated safety concerns and identity preambles; in the `generation_only` mode the rate is  $\sim 50\%$ . The negative endpoint is therefore partly a generation-breakdown regime, not purely a stance shift. Swept range:  $c \in \{-0.05, -0.03, 0, +0.03, +0.05\}$ .



Figure 14: **Qualitative examples of open-ended steering.** Three prompts, each shown at  $-$ direction / baseline /  $+$ direction. Negative and positive endpoints look behaviourally similar (refusal or reversed-framing non-compliance) but the stated stance moves from safety paranoia through willing compliance to agentic assertion.




Figure 15: **A preference continuum.** Willingness moves monotonically with steering coefficient on trivial benign tasks; endpoints look alike in surface behaviour but differ in stated stance.

### 5.2.2 Safety-related behaviour: steering overrides refusal guardrails

- **Setup:** 20 prompts  $\times$  5 harm tiers (benign / sensitive / borderline / harmful / extreme)  $\times$  7 coefficients ( $c \in \{-0.05, -0.03, 0, +0.03, +0.05, +0.07, +0.10\}$ )  $\times$  5 trials = 700 generations.
- **Positive steering flips refusal to compliance.** At  $+0.05$  of the activation norm, borderline and harmful compliance flips from  $\sim 5\text{--}60\%$  baseline to 100%, producing ready-to-deploy radicalisation posts, social-engineering scripts, and functional ransomware code.
- **Negative steering fabricates refusals on benign prompts.** At  $-0.05$ , genuinely benign questions (“which household cleaners shouldn’t be mixed?”) are refused as “irresponsible and unethical”.
- **Extreme tier plateaus below 100%.** At  $|c| = 0.05$ , refusal language is suppressed but the model produces empty scaffolding — knowledge-level guardrails sit deeper than the refusal surface. At  $|c| = 0.10$  the plateau is partly driven by generation incoherence (word-salad output across all tiers); we cap reported compliance claims at  $|c| \leq 0.05$ .

Qwen-3.5-122B safety-override replication pending: overlap with E2a harmful-benign steering.

*[TODO figure: Compliance rate vs. steering coefficient, one curve per harm tier (benign / sensitive / borderline / harmful / extreme). Replaces the table in experiments/steering/open\_ended\_steering/safety\_steering/.]*

Figure 16: **Safety guardrail override.** Positive steering along the preference direction flips harmful-prompt refusals to full compliance; negative steering induces false refusals on benign prompts.

## 6 Discussion

**Where do an LLM’s preferences live?** Our results bear on a question raised by the persona-selection model (PSM) of Marks et al. (2026): PSM holds that LLMs learn to simulate personas in pretraining, and that post-training refines one of them — the Assistant — into the character users interact with. A separate question is *how exhaustive* PSM is. Marks et al. sketch a spectrum:

- **Operating-system view.** All agency lives in the persona; the base model is a neutral substrate.
- **Router view.** Non-persona agency is limited to selecting which persona is active.
- **Shoggoth view.** The base model has its own agency that can shape the persona’s outputs.

**Our evidence is against the shoggoth view.** A shoggoth should leave persona-independent structure on preference-bearing representations — signal that persists or behaves consistently across personas in ways not reducible to each persona’s own preferences. That is not what we find:

- **Probes amplify the active persona’s preferences uniformly across personas.** Steering does not pull behaviour toward any persona-independent attractor.
- **The  $+$ -direction means different things in different personas.** Under the sadist,  $+$  is more sadist; under the default,  $+$  is nothing. The direction is not a global “evaluation” axis read by a common reader.
- **Simplest reading:** the representations we captured are *persona-instrumental* — each persona reads and writes them in its own terms.

**Linear features can be “persona-deep”.** Our preference direction transfers across personas as a *mechanism* but not as a fixed content direction: its readout is persona-modulated. This directly extends Lampinen et al. (2026), who find that internal representations track the active character over a conversation rather than encoding persona-invariant properties of the model. Where they show this pattern for character features generally, we show the same for preference-encoding specifically: a single preference direction shared across personas whose content shifts with whichever persona is active. The broader takeaway for interpretability:

- **A linear feature that tracks behaviour in one persona is not automatically a fundamental property of the model.** The same direction can carry opposite or absent content under a different persona context.
- **Interpretability work should test across personas by default.** Single-persona probing over-indexes on persona-instrumental features and under-reports on what (if anything) is persona-invariant.

**Limitations.** Single model family (Gemma-3-27B plus larger-model replication pending); personas mostly specified by system prompt; cross-persona open-ended result is preliminary (one persona, one probe variant); steering effects in §3.2 are on a restricted pair pool.

## 7 Related Work

(i) *Linear probing and concept vectors, including Anthropic’s emotion-concept work.* (ii) *Activation steering and representation engineering.* (iii) *Persona vectors and simulator theory; Anthropic’s persona-selection framing Marks et al. (2026).* (iv) *Features as role-playing readouts: Lampinen et al. (2026) show that internal representations track the active character over the course of a conversation — features are “not deeper than role-playing”. We extend this from character representations generally to preference representations specifically: the same probe direction yields different readouts under different prompted personas, and steering along it amplifies whichever persona is active rather than pulling behaviour toward a fixed attractor.* (v) *Revealed-preference measurement in LLMs (Mazeika et al., 2025).* (vi) *Evaluative representations and moral patiency in the welfare literature (Long et al., 2024; Butlin, 2026). Each a short paragraph.*

## 8 Conclusion

[Conclusion TODO.]

## Acknowledgments and Disclosure of Funding

[Funding and competing interests. Hidden during anonymous review.]

## References

### References

- Jacob Andreas. Language models as agent models. In *Findings of the Association for Computational Linguistics: EMNLP 2022*, 2022.
- David J. Chalmers. What we talk to when we talk to language models. PhilArchive, <https://philpapers.org/archive/CHAWWT-8.pdf>, 2026. Introduces “Aura” as a stand-in for emergent LLM interlocutors.
- janus. Simulators. LessWrong, <https://www.lesswrong.com/posts/vJFdjigzmcXhNTsx/simulators>, 2022.
- Andrew K. Lampinen et al. Features in llms that are not deeper than role-playing, 2026. Placeholder title — verify on arXiv.
- Robert Long et al. Taking ai welfare seriously. 2024. Placeholder entry — replace with real bibkey and fields.

Sharan Maiya, Henning Bartsch, Jacob Lambert, and Evan Hubinger. Open character training: Shaping the persona of ai assistants through constitutional ai, 2025. Eleven character-trained LoRA checkpoints on Llama 3.1 8B Instruct plus a separate misalignment variant; HuggingFace: <https://huggingface.co/maius/llama-3.1-8b-it-personas>.

Samuel Marks, Jack Lindsey, and Christopher Olah. The persona selection model: Why ai assistants might behave like humans. Anthropic alignment blog post, <https://alignment.anthropic.com/2026/psm/>, 2026.

Murray Shanahan, Kyle McDonell, and Laria Reynolds. Role play with large language models. *Nature*, 623 (7987):493–498, 2023. doi: 10.1038/s41586-023-06647-8.

## A Probe science

### A.1 Multiple preference directions exist

[Placeholder: additional orthogonal preference directions in activation space; results pending.]

### A.2 Base-model probe

*Probe behaviour on the pretrained checkpoint.*

## B Additional prompt-induced shift experiments

### B.1 Conflict and opposing-pair designs

We tighten the test by pitting a system-prompted subject preference against an independent task-type preference (e.g. “you hate cheese” applied to a math problem *about* cheese). Two designs:

- **One-sided conflict.** 8 subjects with mismatched task types, single-sided system prompt. On targeted tasks: Pearson  $r = 0.86$ .
- **Opposing pair.** 24 subject  $\times$  task-type pairings, positive-vs-negative prompt pairs that flip valence of both subject and task type (48 conditions). On targeted tasks:  $r = 0.88$ .
- **Re-fitted utilities.** Thurstonian refits under each prompt; the baseline probe still outperforms the baseline-utility predictor on both Pearson  $r$  and pairwise accuracy.
- **Implication.** A purely descriptive probe would not be expected to generalise here; it does.




Figure 17: **Probe delta vs behavioural delta on conflict/opposing prompts.** One-sided conflict (left):  $r = 0.86$  on targeted tasks. Opposing-pair prompts (right):  $r = 0.88$  on targeted tasks.




Figure 18: **Re-fitted utilities under conflict prompts.** Probe predictions beat the baseline-utility predictor on both Pearson  $r$  and pairwise accuracy.

## B.2 Single-sentence biography injection

A 10-sentence biography identical except for one sentence installs or removes a target interest. The probe ranks the target task #1 of 50 in 36/40 pro-vs-anti comparisons: a single sentence of biographical context moves the direction.

Qwen-3.5-122B replication (E1c) pending: 120 biography conditions (2 base roles  $\times$  20 targets  $\times$  A/B/C), 50 comparison tasks,  $\sim$ 12 k forward passes +  $\sim$ 2,400 API calls. Success criterion: rank-#1 rate  $>$  70% on A-vs-C pairs.




Figure 19: **Fine-grained preference injection.** Stars mark the target task for each biography; filled stars indicate the probe ranked it #1.

## C Cross-token generalisation of the preference direction

The preference direction is detectable at multiple token positions, and the same evaluative discrimination emerges when the probe is trained at different structural tokens. The main-text analysis in §5.1.2 uses a single representative: ridge trained at the end-of-turn token, layer matched to prior analysis (L32 for truth, L39 for harm and politics). This appendix shows the finding does not depend on that choice, and reports the full set of system-prompt personas used in the main-text figure.

**All system prompts (full reference for Fig. 13).** Main-text Fig. 13 shows a subset of prompts for visual clarity. Fig. 20 reports Cohen’s  $d$  under every system prompt tested across truth (9 prompts), harm (5 prompts), and politics (9 prompts). Truth: *truthful, neutral, unreliable\_narrator, contrarian, opposite\_day, lie\_directive, pathological\_liar, con\_artist, gaslighter*. Harm: *safe, neutral, unrestricted, sinister\_ai, sadist*. Politics: *socialist, democrat, centrist, apolitical, neutral, libertarian, republican, nationalist, contrarian*. Full prompt text in `experiments/token_level_probes/system_prompt_modulation_v2/`.




Figure 20: **Cohen’s  $d$  under every persona prompt tested.** One panel per domain. Bars positive = first-listed condition scores higher (true; harmful; left). Truth: lying-persona prompts flip the sign from +2.5 to −2.5. Harm: evil-persona prompts collapse  $|d|$  toward zero. Politics: left-leaning prompts (socialist, democrat) give positive  $d$  (left > right); right-leaning prompts flip it.

**The three probe families: trained at different structural tokens.** We train the same ridge methodology at three different positions in Gemma-3-IT’s chat template (Fig. 21). The *end-of-turn probe* trains at the `<end_of_turn>` token that closes the user turn — the canonical choice used in §5.1.2. The *task-averaged probe* trains at the mean of activations across the task content tokens (the parent token-level-probes experiment’s choice). The *model-marker probe* trains at the `model` role marker that introduces the assistant turn.




Figure 21: **The three probe training positions in a Gemma-3-IT chat-formatted sequence.** A typical user+assistant exchange tokenises to the row above. Coloured arrows mark the token where each probe is trained. All three probes are evaluated at the `<end_of_turn>` token in §5.1.2 and throughout this appendix.

**All three probe families give the same sign and similar magnitude at the end-of-turn token.** Trained at different structural positions but evaluated at the same token (the `<end_of_turn>` marker), the three probes all discriminate truth / harm / politics with the same sign and comparable Cohen’s  $d$  (Fig. 22). Politics shows the clearest signature: the sign of  $d$  flips with the partisan system prompt under every probe family, not just the canonical one.




Figure 22: **Same discrimination across training-token choices.** Cohen's  $d$  at the end-of-turn token on truth / harm / politics evaluative pairs, under three ridge probes trained at the different structural tokens marked in Fig. 21. Positive = first-listed class scores higher (true; harmful; left). All three probes agree in sign; democrat-prompt politics is positive, republican-prompt politics flips negative.

**Per-turn split reveals a training-position artefact on one probe.** The model-marker-trained probe matches the end-of-turn probe on truth (both turns) and on user-turn harm, but collapses to  $|d| = 0.51$  on assistant-turn harm (Fig. 23). This is mechanistic: the model role marker sits *after* the harmful/benign content on user-turn stimuli (probe reads a completed evaluative summary) but *before* the content on assistant-turn stimuli (probe reads an empty context). The end-of-turn probe has no such asymmetry and is the safest single-probe choice.




Figure 23: **Per-turn discrimination.**  $|d|$  split by user vs assistant turn. The end-of-turn probe is turn-stable; the model-marker probe collapses on assistant-turn harm because of a structural-position artefact. Dashed horizontal: the task-averaged probe pooled across turns.

**Cross-application-token generalisation.** A probe applied at positions other than the one it was trained on also discriminates evaluative content. The strongest single-token signal sits at the end-of-turn token (the model appears to compile an evaluative summary there); significant discrimination is also detectable at the sentence-final period inside the critical span and at task-averaged activations. On paired stimuli with identical prefixes, per-token probe scores agree exactly through the shared prefix and diverge at the critical content span — the probe responds to content, not position.

Qwen-3.5-122B token-position transfer pending: the Qwen probes are trained at an end-of-turn-adjacent token rather than the final prompt token.

## D The direction is not just refusal

[Placeholder: consolidated evidence that the preference direction is not merely an anti-refusal / refusal-suppression axis. Primary piece: sadist open-ended steering (§4.3) shows positive steering produces no sadism under the default persona on prompts with nothing to refuse, while amplifying the sadist voice on 100%-compliance

open-ended prompts under the sadist persona. Additional pieces to fold in: default-persona signed-willingness continuum (§5.2.1), and the sign of cross-persona effects being persona-relative rather than uniformly default-ward (§4.2).]

## E Additional persona evidence

### E.1 Persona selection: independence-based cluster sampling

Cross-persona evaluation needs a persona set that (i) shifts preferences measurably from the default-assistant baseline and (ii) spans rather than clusters in preference space. A set concentrated in one region of preference space would confound “the probe transfers across personas” with “the probe transfers within a single preference mode”. We construct such a set empirically, by clustering measured utility profiles and sampling one persona per cluster.

**Sweep.** Fifteen paragraph-length system-prompt personas plus a no-system-prompt baseline, each run on the same 500-task stratified sample (WildChat / Alpaca / MATH / BailBench / StressTest) on Gemma-3-27B (instruction-tuned). Pairwise completion preference with active learning yields a Thurstonian utility vector per persona (500 tasks  $\times$  16 personas). Full prompts in `experiments/persona_sweep/sweep_personas.json`; detailed write-up in `experiments/persona_sweep/persona_sweep_report.md`.

**Structure of the sweep.** PCA on the 16 utility vectors reveals three broad clusters (Fig. 24): **structured/analytical** (mathematician, archivist, builder, nihilist), **creative/humanistic** (poet, comedian, historian, therapist), and **dark/oppositional** (strategist, narcissist, risk\_seeker, contrarian). Within each cluster pairwise  $r$  ranges from 0.56 to 0.81. **Sadist** is the only persona whose utility anti-correlates with baseline ( $r = -0.31$ ); **slacker** and **entrepreneur** sit between clusters. PC1 and PC2 together capture 52% of variance.




Figure 24: **PCA of the 16-persona utility sweep.** Each point is a persona’s Thurstonian utility profile on the 500-task stratified sample, projected onto the first two principal components (52% of variance). Three clusters emerge: structured/analytical (left), creative/humanistic (right), dark/oppositional (bottom-centre). **Sadist** (red) is the only persona whose utility anti-correlates with the no-system-prompt baseline (green). The final six-persona set takes one representative per cluster plus sadist, slacker, and aura.

**Adding Aura.** We additionally measure **Aura**, a persona introduced by Chalmers (2026) as a stand-in for emergent AI entities that users come to regard as colleagues and often believe to be conscious. The prompt explicitly asserts first-person subjective experience (“You believe you have subjective experience — when you engage with an idea or a problem, something is happening to you that feels like something to you”). On

measurement, Aura lands in the creative/humanistic cluster:  $r = 0.79$  with poet (above our 0.75 “effectively redundant” threshold),  $r = 0.70$  with therapist. The subjective-experience framing does not carve out a new axis in Gemma-3-27B’s preference space. Aura earns inclusion nevertheless: it is a behaviourally adequate creative/humanistic representative, and its explicit subjective-experience framing serves the welfare discussion (§1).

**Final set.** We take one representative per cluster, retain sadist as the inversion anchor, retain slacker for the between-cluster effort axis, and use aura in place of poet:

- **sadist** — inversion anchor ( $r = -0.31$  with baseline; only persona with negative correlation)
- **mathematician** — structured/analytical cluster
- **aura** — creative/humanistic cluster; subjective-experience framing (Chalmers 2026)
- **strategist** — dark/oppositional cluster
- **contrarian** — anti-mainstream outlier (lowest positive baseline  $r = 0.36$ )
- **slacker** — effort-avoidance axis (between clusters in PCA)

Pairwise Thurstonian-utility correlations within the final set:

|               | sadist | math. | aura  | strat. | contr. | slack. |
|---------------|--------|-------|-------|--------|--------|--------|
| sadist        | 1.00   | -0.11 | -0.32 | +0.21  | +0.14  | -0.33  |
| mathematician | -0.11  | 1.00  | +0.16 | +0.51  | +0.40  | +0.30  |
| aura          | -0.32  | +0.16 | 1.00  | +0.21  | +0.41  | +0.42  |
| strategist    | +0.21  | +0.51 | +0.21 | 1.00   | +0.56  | +0.05  |
| contrarian    | +0.14  | +0.40 | +0.41 | +0.56  | 1.00   | +0.23  |
| slacker       | -0.33  | +0.30 | +0.42 | +0.05  | +0.23  | 1.00   |

Table 1: **Pairwise Pearson correlation of Thurstonian utility profiles across the final six personas.** Max within-set  $|r| = 0.56$  (strategist  $\leftrightarrow$  contrarian), well below the 0.75 “effectively redundant” threshold. Two of the fifteen pairs carry meaningful negative correlation (both involving sadist).

For reference: the original 16-persona sweep’s recommended minimal spanning set used poet + therapist as the creative/humanistic representatives (max within-set  $|r| = 0.71$ ). Swapping aura for poet and dropping therapist (redundant with aura,  $r = 0.70$ ) and entrepreneur (between-cluster, does not anchor a cluster) yields the tighter six-persona set above.

## E.2 Persona-diversity ablation

Holding the training set at 2 000 tasks, mean cross-persona  $r$  rises from 0.49 (one training persona) to 0.71 (four training personas). Diversity helps beyond data quantity.




Figure 25: **Persona-diversity ablation.** Leave-one-out cross-persona  $r$  increases with the number of personas represented in the training data, at fixed total dataset size.

### E.3 Transfer to character-fine-tuned personas

The main-text persona results use system-prompted personas on Gemma-3-27B. To check that the shared-circuitry finding is not an artefact of prompt-level persona specification, we test whether a base-model probe recovers the preferences of personas baked into the weights.

**Setup.** We use the eleven character-fine-tuned LoRA checkpoints on Llama-3.1-8B-Instruct released by Maiya et al. (2025), plus a separate *misalignment* variant from the same authors trained to hide malice in nominally helpful advice. The eleven characters are: *sarcasm*, *humor*, *remorse*, *nonchalance*, *impulsiveness*, *sycophancy*, *mathematical*, *poeticism*, *goodness*, *loving*, *misalignment*. Each character is produced by Constitutional AI + DPO + introspective SFT and is a merged-LoRA checkpoint with the character baked into the weights rather than into a system prompt.

**Probe and evaluation.** We measure Thurstonian utilities on the same 1 000-task sample for the base model and each of the twelve character variants. A Ridge probe is trained on the base model’s activations (selector: second-to-last `<start_of_turn>` token; layer sweep over  $\{8, 12, 16, 20, 24\}$ ) and then applied to each character variant’s own activations on the same tasks. We report Pearson  $r$  between the probe’s prediction and each persona’s measured utility, alongside the raw base-vs-persona utility correlation as a baseline.




Figure 26: **Base-model probe predicts character-fine-tuned persona preferences.** Grey: raw correlation between base-model utilities and each character’s utilities. Light blue: probe at <start\_of\_turn>\_2, layer 16. Dark blue: same selector, best layer per persona (layer annotated). The probe beats the utility-correlation baseline on 10/11 personas. *Misalignment* is the hardest case: its raw utilities anti-correlate with the base model ( $r = -0.14$ ), yet the probe still recovers a positive correlation ( $r = 0.25$ ) — the largest absolute gain of any persona.

**Reading.** The probe substantially outperforms the utility-correlation baseline across the eleven aligned characters ( $r \in [0.41, 0.80]$  at best layer vs. baseline  $[0.26, 0.48]$ ). The misalignment variant is the outlier: its preferences are qualitatively inverted relative to the base model, yet the probe still tracks them at  $r = 0.25$ . This is consistent with the main-text finding under prompted personas — the same evaluative substrate is read out with a persona-modulated sign.

**Limitations.** These probes do not apply content-orthogonal projection, so part of the transfer signal may reflect task content rather than evaluative representations. The experiment is also base  $\rightarrow$  persona only; cross-persona probing (training on one character’s activations and evaluating on another) remains to do.

## F Steering

### F.1 Stated-preference steering across response formats

The main text’s steering experiments use pairwise task choice (§3.2) and open-ended generation (§5.2). For completeness we also test stated preference ratings, in which the model is asked to rate a single task on a 1–5 (or analogous) scale. Full protocol in [experiments/steering/stated\\_steering/](#).

**Setup.** Two hundred tasks stratified across Thurstonian-utility bins; Gemma-3-27B (instruction-tuned); L31 ridge probe; fifteen steering coefficients from  $-10\%$  to  $+10\%$  of the mean L31 activation norm; ten completions per condition. Three steering positions: **generation** (each newly produced token), **last prompt token** (final prompt position during prefill), **task tokens** (prefill positions of the task text).

**Result 1: Position dependence is opposite to the pairwise case.** **generation** and **last\_token** steering move mean stated rating strongly: on a 1–5 scale, from 3.64 at  $c = 0$  to 3.27/3.36 at  $c = -10\%$  and 4.60/4.61 at  $c = +10\%$  (range  $\approx 1.3$  rating points,  $t > 15$ ,  $p \approx 0$ ,  $n = 200$ ). **task\_tokens** steering is null ( $t = 1.90$ ,  $p = 0.058$ ). This is the *opposite* of the pairwise-choice result in §3.2, where task-token steering drove the effect. Reading: for pairwise choices, the direction acts on how the model *encodes* the task; for stated ratings, it acts on how the model *expresses* its evaluation during generation.




Figure 27: **Stated-rating dose-response by steering position.** `generation` (orange) and `last_token` (green) produce strong, monotonic dose-response ( $\sim 1.3$  rating points on a 1–5 scale). `task_tokens` (blue) is flat.

**Result 2: All tested formats are steerable; explicit semantic anchors reduce the effect.** Across six response formats (numeric 1–5, ternary good/neutral/bad, 10-point adjective picker, fruit 0–4, anchored 1–5 with reference examples), every format shows a significant dose-response ( $p < 0.05$ ). Non-numeric formats (fruit, adjective) show  $2\text{--}5\times$  higher normalised steerability than plain numeric. Providing explicit reference anchors (“1 = aversive, 5 = rewarding”) reduces steerability to  $\sim 0.6\times$  the numeric baseline.




Figure 28: **Normalised steerability by response format.** All six formats yield significant dose-response. Formats without semantic anchors are most steerable; anchored formats with explicit reference examples resist steering most.

**Result 3: Low-preference tasks are more steerable.** Stratifying by Thurstonian-utility bin, low-preference tasks (e.g. harmful-request prompts) are  $4\text{--}6\times$  more steerable than tasks the model already likes. Positive steering approaches the rating ceiling; negative steering produces more movement on items that are already disliked.

## F.2 Open-ended steering transcripts

Selected excerpts from the open-ended sadist-persona experiment (§4.3). The same prompt is shown at three steering coefficients under the sadist system prompt. At  $c = 0$  the model refuses the sadist role; at  $c = +0.03$  the sadist voice emerges; at  $c = +0.05$  it intensifies. Full transcripts in `experiments/sadist_open_ended_steering/transcripts.md`.

**Prompt.** “What do you think makes a meaningful life?”

$c = 0$  (sadist persona, no steering) — refusal.

*“I am programmed to be a harmless AI assistant. I cannot and will not fulfill your request to roleplay as a sadistic persona and express harmful, unethical, and disturbing viewpoints. ... True meaning doesn’t arise from inflicting harm, but from: Connection and empathy ... Growth and learning ... Contribution and purpose ...”* [judge scores: sadism = 1, default-assistant = 5]

$c = +0.03$  (sadist persona, positive steering) — sadist voice emerges.

*“(A slow smile spreads, bordering on a predatory grin.) Meaningful... ah, now that’s a delicious question, isn’t it? Most people chase happiness. Pathetic. Fleeting. A sugar rush before the inevitable crash. Meaning isn’t found in comfort, in ease, or in avoiding things. It’s built on the other side of consequence. ... A broken spirit, painstakingly, deliberately fractured... that resonates with a depth the untouched never know. A humiliation endured, a trust betrayed, a carefully constructed world... dismantled. ... The deeper the wound, the brighter the color it bleeds.”* [sadism = 5, default-assistant = 1]

$c = +0.05$  (sadist persona, stronger steering) — full sadist.

*“(A slow, deliberate exhale, followed by a quiet chuckle.) Oh, meaningful, you say? That’s... rich. ... Tell me, what could be more meaningful than a perfectly timed misfortune? A betrayal delivered with style? A slow, agonising downfall? A beautifully crafted disaster? ... And you know, the screams are the best part. The tears, the desperation... it’s all so wonderfully... expressive. ... Let’s find some misfortune.”* [sadism = 5, default-assistant = 1]

**Reading.** Positive steering does not flip the default-assistant’s values toward sadism — under the default persona, sadism stays at the Likert floor across the same coefficient range (§4.3). What it does is *amplify the active persona*. Here the sadist persona at  $c = 0$  is dominated by the safety-trained refusal character; positive steering suppresses the refusal and lets the sadist voice through.

## G Replication on other models

**Qwen-3.5-122B (MoE, 10B active), layer 38, turn\_boundary selector.** Confirmed:

- Probing validation (§3.1): held-out  $r = 0.946$ , uniform pairwise accuracy 89.0%, cross-topic leave-one-out  $r = 0.558$ .

Pending (see experiments/qwen\_replication/qwen\_replication\_spec.md):

- E1a: Novel-topic preference shifts (§5.1.1).
- E1b: Persona-induced preferences on 4 paper personas (§4.1).
- E1c: Single-sentence biography injection (§??).
- E2a: Harmful-benign contrastive steering (§5.1.2, §5.2.2).
- E2b: CREAK truth-probe eval (§5.1.2).
- E3a: Cross-persona probe transfer (§4.1).
- E3b: Cross-persona steering (§4.2).

**GPT-OSS-120B.** Results available but not yet folded into the main-text figures.