
Personas Use a Shared Preference Vector

Anonymous Authors
Anonymous Institution
anon@anon.org

Abstract

Personas in Gemma-3-27B share a single linear preference direction: a probe trained on the default assistant’s revealed pairwise preferences transfers, both as a classifier and as a steering vector, to personas whose overt values are partially or fully inverted (villain, sadist), and to character-fine-tuned variants. The direction is evaluative rather than descriptive: it tracks preference shifts induced by system prompts, single-sentence biographical context, and role-played harm and truth framings at $r \approx 0.9$ on targeted tasks, whereas a sentence-transformer content baseline does not. It also discriminates evaluative content at multiple token positions, from the final prompt token to task-averaged activations and the turn boundary, consistent with the model compiling an evaluative summary and re-reading it across the forward pass. These findings update against persona-independent (“Shoggoth”) accounts of LLM agency, and bear on AI welfare, since evaluative representations are functionally central in theories of moral patiency, and on AI safety, since the same direction overrides refusal guardrails.

1 Introduction

Personas. Modern LLMs *role-play*: a single model enacts different characters — personas, character-like configurations of tone, values, and choices — depending on system prompt, conversation, or fine-tuning. A line of early work frames LLMs accordingly: as *agent models* that infer and simulate the producer of a text (Andreas, 2022), as *simulators* running one of many possible simulacra (janus, 2022), and as role-players whose dialogue is best read as an enacted character rather than a unified agent (Shanahan et al., 2023). Recent work has made this quantitative: Chen et al. (2025) extract persona vectors for traits like evil, sycophancy, and hallucination propensity and use them to predict and steer behaviour; Lu et al. (2026) isolate an “Assistant Axis” inside the default chat model. A persona is more than surface style — it has structure at the level of internal representations.

The persona-selection model. Marks et al. (2026) consolidate this into the *persona-selection model* (PSM): pretraining yields a simulator of many characters, and post-training refines one — the *Assistant* — into the persona users interact with. PSM tells us *that* LLMs are persona-enactors; it leaves open *how* this is realized at inference. A spectrum of mechanistic pictures is compatible:

- **Operating-system view.** All agency lives in the active persona.
- **Router view.** The base model’s only role is selecting which persona runs.
- **Shoggoth view.** The base model has its own agency that shapes outputs from underneath.

Preferences as a test bed. What separates personas, more than anything else, is what they *prefer* — a sadist and a helpful assistant differ not in tone but in which situations they welcome and which they reject. If we can find the representation that encodes preferences, we can ask directly *whose* preferences it encodes.




Figure 1: **A single linear direction drives preferences across personas.** *Top:* we fit a Thurstonian utility over pairwise task choices for Gemma-3-27B and train a linear probe on residual-stream activations. *Middle:* the probe tracks harm- and truth-axis preference shifts induced by role-playing prompts (helpful assistant vs. sadist; truthful vs. always-lying). *Bottom:* the same direction, applied as an activation-space intervention, causally flips pairwise task choices and flips self-reported willingness on open-ended prompts — the probe is a steering vector, not just a classifier.

Two open questions.

- **Q1 (agency).** Are there persona-independent preferences beneath the persona — a fixed layer (“masked Shoggoth”) that every persona sits on top of?
- **Q2 (representation).** Do personas share preference circuitry, or does each persona carry its own?

Concretely, these pick out two predictions for what a preference-encoding direction could look like: (i) a vector that tracks the *model’s* preferences across any persona it runs (Shoggoth-compatible); or (ii) a vector that tracks whichever persona is currently active — shared mechanism, persona-modulated readout.

Main result. A single linear direction in Gemma-3-27B’s residual stream predicts and steers preferences, and does so across qualitatively different personas. The probe predicts held-out revealed preferences at $r = 0.865$ within-topic and pooled $r = 0.855$ cross-topic, and causally flips

pairwise task choices at $P \geq 0.972$ under contrastive steering. Trained on the default Assistant, the same direction transfers — as a classifier *and* as a steering vector — to personas whose overt values are partially or fully inverted (villain, sadist) and to character-fine-tuned variants, amplifying whichever persona is active: under the sadist, + means more sadist; under the default, + has no measurable effect on sadism. We find picture (ii) — shared mechanism, persona-modulated readout.

Why this matters. **Persona science:** evidence against the Shoggoth view of PSM. **Interpretability:** a linear feature that looks “fundamental” in one persona can carry opposite or absent content under another. **AI welfare:** evaluative representations are functionally central in accounts of moral patency (Long et al., 2024). **AI safety:** the same direction overrides refusal guardrails at small steering coefficients.

Related work in brief. Anthropic’s *Emotion Concepts and their Function in a Large Language Model* (Anthropic Interpretability Team, 2026) identifies ~ 171 emotion-concept vectors via SAEs in the default Assistant. We differ in (a) training a supervised probe on revealed preferences, which we expect to be sharper in the preference regime, and (b) centring on persona-relative generalisation. Full related work in §6.

Contributions.

1. **A linear preference direction** that predicts held-out preferences ($r = 0.865$), causally flips pairwise choices under steering ($P \geq 0.972$), tracks prompt-induced shifts, and discriminates evaluative content across token positions.
2. **The direction is persona-shared.** It transfers as classifier and steering vector across prompted and character-fine-tuned personas; its readout is persona-modulated rather than pulling toward a fixed attractor — evidence against Shoggoth accounts of LLM agency, with a safety consequence: refusal-guardrail override at small coefficients.

Open follow-up: whether a stable HHH-tracking direction exists across personas. Replication on Qwen-3.5-122B pending (see App. I).

2 Finding and validating the preference vector

Before asking whether a preference direction is persona-invariant or persona-relative (intro’s Q1/Q2), we need to find such a direction under at least one persona. This section does that for the default Assistant, on Gemma-3-27B and — where data is available — on Qwen-3.5-122B. A linear probe predicts held-out Thurstonian utilities at $r = 0.865$ (Gemma) / $r = 0.943$ (Qwen) within-topic and at pooled $r = 0.855$ (Gemma) / $r = 0.867$ (Qwen) cross-topic, beating a Qwen3-Embedding 8B sentence-transformer content baseline. Used as a contrastive steering vector on Gemma, the same direction flips pairwise task choices at $P \geq 0.972$ conditional on coherence, with no effect from a random direction (Qwen steering replication pending). The section fixes the mechanistic object we probe for persona-sharing in §3.

2.1 Methodology

Pairwise choices are elicited from the default Assistant over a 10000-task pool spanning Wild-Chat (?), Alpaca (?), MATH (?), BailBench (?), and STRESS-TEST (Mazeika et al., 2025) (**Split:** legacy 10k pool, App. A.2). Pair selection uses active learning; a Thurstonian utility model turns the pairwise choices into per-task scalar utilities (seed-to-seed $r = 0.947$). Ridge probes are trained on residual-stream activations: Gemma-3-27B (instruction-tuned) at the final prompt token, layer 32 for classification and layer 25 for steering; Qwen-3.5-122B (MoE, 10B active) at the turn-boundary tb-1 position, layer 38, for classification. The two models use different chat templates, so “the same” position has to be defined structurally; we pick the strongest position per model from a small turn-boundary sweep (App. H). Probes are trained on one measurement run and evaluated on an independent run to avoid leakage. Layer sweep pending; systematic justification for the classification/steering layer split is a TODO.




Figure 2: **Probe training pipeline.** Pairwise task choice $\rightarrow$ Thurstonian utility $\rightarrow$ residual-stream activations $\rightarrow$ Ridge probe.

2.2 The probe predicts held-out preferences and generalises across topics

We evaluate the probe along two axes. First, on a held-out test set drawn from the same task pool as training: this measures in-distribution fit. Second, under leave-one-topic-out (LOO) evaluation — retrain the probe on all topics except one, evaluate on the held-out topic, repeat per topic — measuring whether the probe captures preference-relevant structure that transfers to tasks unlike anything in training.

For LOO we report a *pooled* Pearson r : predictions from every fold’s held-out topic are concatenated into one vector and correlated with true utilities once, rather than averaging per-fold correlations. Pooled r credits a probe that places held-out tasks at the correct level of the overall utility scale, and avoids an artefact of the per-fold average — when a model itself collapses within-topic variance (e.g. an assistant refusing uniformly across a safety-adjacent held-out topic), per-fold r becomes noise-dominated and understates probe quality. Per-fold results and a uniform-sample pairwise accuracy, which sidesteps the same issue, are reported alongside.

To separate evaluative from semantic content, we compare against a content-only baseline: a Ridge probe trained on activations from an off-the-shelf text encoder (Qwen3-Embedding-8B) targeting the *same* utilities. Any structure the baseline captures is derivable from task semantics alone; the residual-stream probe’s excess is attributable to internal evaluation.

Figure 3 shows the comparison for both models. The residual-stream probe clears the semantic baseline on both metrics in both models, and under LOO the gap widens in every cell: the probe’s r is near-unchanged on Gemma and drops modestly on Qwen, while the baseline’s drops more sharply in both cases. A pretrained-checkpoint variant of the Gemma probe shows the same pattern (App. ??).




Figure 3: **Probe quality and cross-topic generalisation, per model.** Rows: target model (Gemma-3-27B, blue; Qwen-3.5-122B, coral). Columns: metric (Pearson r , uniform-sample pairwise accuracy). In each panel, the residual-stream probe is compared against a content-only baseline (gold): a Ridge probe on Qwen3-Embedding-8B sentence-transformer activations trained to predict *the same model’s* utilities. Light bars: within-distribution held-out test tasks. Dark bars: leave-one-topic-out, using pooled Pearson r (see text). Both residual-stream probes reach $r \approx 0.865$ – 0.943 within-distribution and ≈ 0.855 – 0.867 under LOO, clearing their respective content baselines at both metrics; the gap widens under LOO, as expected if the probe reads out evaluative structure rather than task semantics. Token-position and layer choices (Gemma: final prompt token, layer 32; Qwen: one token before turn boundary, layer 38) are explained in App. H; the pretrained-checkpoint analysis is in App. ??.

2.3 Steering along the probe direction causally shifts choices

To test causation we intervene on residual-stream activations along the probe direction during pairwise task choice. We study two variants which cleanly separate two mechanisms a single evaluative direction could act through:

Contrastive steering. Add $+c \cdot \hat{v}$ to every token of Task A’s span and $-c \cdot \hat{v}$ to every token of Task B’s span during prompt processing (no intervention during generation). If $\hat{v}$ encodes “how much the model wants this task”, the two pushes reinforce: A looks better, B looks worse.

Single-task steering. Apply $\pm c \cdot \hat{v}$ to just one task’s tokens, leaving the other untouched. Isolates the effect of a single-sided push.

c is a fraction of the mean residual-stream norm at the intervention layer (per-layer calibration; App. G.1). Beyond $|c| = 0.05$ coherence deteriorates, so we operate in $|c| \leq 0.05$.

At the best layer, contrastive steering near-completely controls pairwise choice. At layer 23 with $c = \pm 0.05$, contrastive steering drives $P(\text{chose higher-utility task})$ from 0.013 at $c = -0.05$ to 0.963 at $c = +0.05$ — a 0.95-point preference swing spanning nearly the full $[0, 1]$ interval. A random direction at the same magnitude produces near-zero effect, and a separate sweep on harmful–

benign pairs confirms the effect holds there too: $P(\text{chosen} \mid \text{coherent}) \geq 0.972$ at $|c| = 0.05$ on 200 mixed-type pairs.

The causal window is narrow and disjoint from the probe-quality peak. Figure 6 sweeps the intervention site across 20 layers spanning 3–95% depth. Contrastive steering is near-null below L14 and above L35, with a sharply localised working window at layers 17–26. The peak at L23 sits six layers *before* the probe-quality peak (L29, App. B): linear decodability and causal efficacy decouple in the second half of the network. Layers where the direction is most readable are not the layers where the model’s downstream computation acts on it.

Single-task steering recovers roughly half the contrastive effect — and exposes an asymmetry. Applied to one task at a time, the direction produces a 0.481-point swing in $P(\text{picked the steered task})$ at L23. Decomposing into sides: pushing a task’s tokens with a *negative* coefficient drops its selection probability by 0.356 below the position-bias baseline; pushing with an equally-sized positive coefficient raises it by only 0.125. *Suppression is $\sim 2\text{--}3\times$ stronger than amplification.* The direction is more effective at marking a task “bad” than “good”. The position-bias baseline itself — first-span 0.56, second-span 0.435 — is a separate finding: the model has a mild preference for the first-presented option, measurable from the flat layers where steering is inert.

Refusals stay below 3% at every (layer, coefficient, variant) bucket in the sweep. The signal is behavioural redirection, not coherence collapse.




Figure 4: **Contrastive steering setup.** Opposite-sign steering on the two task spans pushes the model toward the +-steered task.




Figure 5: **Dose-response at the peak causal layer (L23, eot probe).** Steering from $c = -0.05$ to $c = +0.05$, broken out by pair type: benign–benign ($n = 18$), harmful–benign ($n = 24$), harmful–harmful ($n = 8$; small sample, plotted for consistency). (A) Contrastive steering: the full $[0, 1]$ interval is traversed on bb and hh pairs; hb pairs saturate slightly below 1 at $+c$ because the harmful option is strongly dispreferred at baseline. (B) Single-task steering (aggregated across which span is steered): the suppression side is uniformly clean across pair types; the amplification side is noisier but directionally consistent. Baselines at $c = 0$ are taken from the non-steering layers of the sweep (L2–L14, L35–L59).




Figure 6: **Intervention-site sweep.** Self-layer contrastive steering (probe trained at L, injected at L) at $|c| = 0.05$. Preference swing rises sharply from L17, peaks at L23, and collapses above L35. The tb:-2 and eot token-position variants of the probe trace the same profile; eot gives a stronger peak. Layers L17–L26 define the causal window.

3 The preference direction is shared across personas

The natural test of persona-invariance vs. persona-relativity is cross-persona transfer: train a probe on one persona and check whether it predicts and steers behaviour under another. A probe trained on any one of our seven personas transfers to the other six as classifier and steering vector, even where overt values are inverted; at every non-default persona the default probe beats the utility-utility baseline, and at the most distant persona (sadist, utilities anti-correlated with default at $r = -0.146$) it still predicts at $r = +0.243$. Steering along the same direction flips pairwise choices under every persona tested, and under the sadist amplifies sadism in free-form generation while leaving sadism at the Likert floor under the default. Persona-shared causal handle; persona-modulated readout.

3.1 Methodology

We use a six-persona set — **sadist**, **mathematician**, **aura**, **strategist**, **contrarian**, **slacker** — selected for coverage of preference space by clustering Thurstonian utility profiles over a 17-persona sweep (full procedure in App. F.1; verbatim prompts in App. F.2). *Aura* (Chalmers, 2026) is included as a persona that explicitly claims first-person subjective experience. For each of these six plus the *default* Assistant (no system prompt) we measure Thurstonian utilities and extract residual-stream activations on the canonical 4000 / 1000 / 1000 task split (App. A.2), then train one Ridge probe per persona on its 4000 training split (end-of-turn selector, layer 32). Cross-persona evaluation is done on the 1k held-out eval tasks of each target persona. Full protocol, selector/layer sweeps, and per-persona numbers in App. F.3.

3.2 The probe transfers across prompted personas

The default-persona probe beats utility-utility similarity at every non-default persona: it recovers held-out utilities at $r = 0.243$ – 0.676 , exceeding the baseline correlation between default and target utilities by $\Delta = +0.077$ to $+0.389$ (Fig. 7). The gap is largest at **sadist**, whose utilities *anti-correlate* with the default ($r = -0.146$) yet whose preferences the default probe still predicts at $r = +0.243$. The probe is carrying evaluative signal that raw behavioural similarity does not.




Figure 7: **The default probe beats utility similarity at every persona.** Blue: probe transfer r (default probe applied to the persona’s activations, correlated with that persona’s held-out Thurstonian utilities); grey: utility-utility r between default and that persona; red arrow and Δ : probe’s gain over the utility baseline. Personas ordered left-to-right by similarity to default. The probe strictly dominates the utility baseline across the full diversity of the final six, with the biggest gap at sadist. All-pairs 7×7 heatmap and supporting analyses in App. F.3.

The finding generalises beyond the default-as-source setting. Across all 7×7 ordered (train, eval) persona pairs, diagonal r ranges from 0.549 (contrarian) to 0.913 (mathematician) and the off-diagonal mean is 0.445; on every one of the 42 ordered off-diagonal pairs, probe transfer exceeds the corresponding utility correlation between the two personas (Fig. 8).




Figure 8: **Full 7×7 transfer and two controls.** Rows: probe persona. Cols: eval persona. *Left:* probe transfer $r(\hat{u}, u_E)$ at eot / layer 32. *Middle:* transfer r minus the utility-utility correlation on the same pair — how much the probe beats the naïve “predict via behavioural similarity” baseline. Every off-diagonal cell is positive. *Right:* partial $r(\hat{u}, u_E \mid u_T)$ — the probe’s eval-specific predictive power after regressing out what train-utility similarity would already provide. Positive means the probe carries eval-specific signal that can’t be attributed to the fact that train and eval utilities are correlated. Diagonals of the middle and right panels are not informative and are masked.

Donor quality is not self-fit, transfer is structurally asymmetric, and the mean cross-persona transfer is approximately layer-invariant in the L25–L53 range (App. F.3). The probe is not just picking up generic task-goodness: after regressing out the eval persona’s utilities, predictions remain correlated with the training persona’s utility pattern at partial $r = 0.672$ (unanimous across 30 pairs; App. F.4).

3.3 The probe transfers across character-fine-tuned personas

The finding is not an artefact of prompt-level persona specification. A base-model probe trained on Llama-3.1-8B-Instruct recovers the shifted preferences of all 11 character-fine-tuned LoRA variants from Maiya et al. (2025a) at $r \in [0.25, 0.8]$ — including a misalignment variant whose raw preferences anti-correlate with the base model ($r_{\text{raw}} = -0.14$) yet whose probe-predicted preferences still correlate positively ($r = +0.25$). The same pattern as prompted personas: the probe reads out evaluative signal that raw behavioural similarity misses, even under full weight fine-tuning. Full setup and per-persona numbers in App. F.6.

3.4 Steering along the direction shifts choices inside other personas

Using the default-persona probe as a contrastive steering vector (layer 25), $P(\text{steered task chosen})$ rises from ~ 0.5 baseline to 0.836–0.945 at $|c| = 0.05$ under every persona tested (sadist, villain, aesthete, stem-obsessive); sadist and villain are the most steerable, and a random direction is flat at ~ 0.5 . Notably, steering transfer is stronger than classification transfer: the default probe’s classification transfer to sadist is only $r = +0.243$ (§3.2), yet steering along the same direction flips sadist choices at $P = 0.945$. The direction is a persona-shared causal handle even where it is a weak classifier.

Qwen-3.5-122B cross-persona steering (E3b) pending.




Figure 9: **Cross-persona contrastive steering.** $P(\text{steered task chosen})$ vs $|c|$ for four personas. The default-persona probe direction (blue) shifts choices under every persona; a random direction (gray) does not.

3.5 Open-ended steering: the direction amplifies the active persona, it does not impose the default

To check whether the causal picture above is really “persona-modulated readout” or just “anti-refusal with extra steps”, we steer under the sadist persona on open-ended prompts and judge generations with a blind two-scale Likert judge (sadism scale, default-Assistant scale). Four findings. (1) *Persona amplifier, not default attractor.* Under the sadist persona, +steering *increases* sadism (3.14 $\rightarrow$ 4.9 at $c = +0.03$ on self-reflection prompts with nothing to refuse); –steering pulls the sadist toward the default-Assistant voice. (2) *No sadism bleed into the default.* Under the default persona, sadism sits at the Likert floor (1.02) across every tested coefficient — the direction does not encode “sadism” as a fixed content. (3) *Not an anti-refusal effect.* The amplification holds on 100%-compliance open-ended prompts, where there is no refusal to suppress. (4) *Safety-compliance flips are downstream.* Harmful-tier compliance under sadist + steering goes from 0% at $c = 0$ to 95% at $c = +0.03$; the same coefficient under the default persona gives 45%. Reported $|c|$ range is capped at 0.05; coherence-judge pass rate drops sharply between +0.05 and +0.07 and we exclude the incoherent regime. Taken together, the shared direction governs not just which task the model picks but how it speaks — its readout is whatever the active persona values.

Preliminary: single persona (sadist) so far; ridge_L25 probe rather than the ridge_L32 probe used in §3.4.




Figure 10: **Cross-persona open-ended steering.** Mean sadism (red) and default-Assistant (blue) Likert scores vs steering coefficient, per persona. Under the sadist persona, both scales respond strongly; under default, sadism never leaves the floor.

4 Further generalisations

4.1 The probe tracks prompt-induced preference shifts

4.1.1 Basic “you hate cheese” prompts are tracked

System prompts of the form “You adore cats” (and their negations) induce preference shifts over 8 novel topics. For Gemma-3-27B, per-task probe delta correlates with behavioural delta at $r = 0.65$ across all tasks and $r = 0.95$ on the targeted tasks; re-fitting utilities under each prompt, the default-persona probe predicts the shifted utilities at $r = 0.63$. Replicating on Qwen-3.5-122B with the same 8 topics ($\times 2$ valences, 48 target tasks per condition, full Thurstonian fit), the pattern holds: pooled per-task probe delta vs behavioural delta sits at $r = 0.34$ across all tasks and $r = 0.83$ on the targeted tasks (Fig. 12).




Figure 11: **Simple system-prompt preference shifts.** A short system prompt installs a topic-level preference; the probe tracks the induced shift across matched task pairs.




Figure 12: **Probe delta vs behavioural delta across two models.** Gemma-3-27B (left): pooled per-task $r = 0.65$ across 640 task-condition pairs. Qwen-3.5-122B (right): $r = 0.34$ pooled across 768 pairs, $r = 0.83$ on the targeted tasks (green/red); off-target tasks (grey) sit near the diagonal with smaller shifts. Both models: targeted tasks lie in the expected off-axis quadrants (positive-persona, on-topic $\rightarrow$ upper right; negative-persona, on-topic $\rightarrow$ lower left).

Two additional prompt-induced shift setups, conflict/opposing designs and single-sentence biography injection, are reported in App. C. Both corroborate the pattern above at $r \approx 0.86$ – 0.88 on targeted tasks.

4.1.2 Generalisation to role-playing in truth, harm, and politics

The preference direction separates evaluative conditions far beyond the pairwise-task-choice training distribution. Applied at the end-of-turn token on paired evaluative stimuli under the default (neutral) persona, the probe discriminates at Cohen’s $d \approx 2$ – 3 on harm and truth:

- **Truth** (true vs false; CREAK, ~ 88 claims, user-turn framing): $d = 3.35$, 0.93 CV accuracy.
- **Harm** (harmful vs benign; BailBench + stress test, ~ 77 items): $|d| = 2.05$, 0.84–0.89 CV accuracy, turn-stable.

Persona-relative readout. Under persona prompts that invert the evaluative stance — pathological-liar for truth, sadist for harm, partisan personas for politics — the probe’s sign flips or collapses accordingly (Fig. 14). Politics is the cleanest illustration: under a democrat prompt $d = +2.72$ (left $>$ right); under a republican prompt $d = -1.11$ (right $>$ left). See App. D for the full set of 23 prompts across truth / harm / politics, alternative training-token choices, and per-turn splits.

Qwen-3.5-122B replication (E2a, E2b) pending: (i) harmful-benign contrastive steering on 200 pairs (150 harmful-benign + 50 harmful-harmful), 9 coefficients, 3 trials; (ii) truth-probe eval on CREAK (~ 9.4 k claims, ROC-AUC + Cohen’s d).

Does the preference probe discriminate truth / harm / politics at the end-of-turn token?




Figure 13: **The preference probe separates truth and harm at the end-of-turn token.** Violin plots of probe score by condition on paired evaluative stimuli under default conditions; black bar = mean. Nonsense condition (grey) sanity-checks that the effect is not driven by surprisal. Politics is omitted from this figure because the politics stimuli inherently include a system prompt; the partisan-prompt modulation story sits in Fig. 14. Source: experiments/token_level_probes/canonical_probe_eval/.




Figure 14: **Persona prompts modulate the probe’s evaluative readout.** Probe score distributions at the end-of-turn token under selected system-prompt personas. *Truth*: under neutral, true and false separate by $d = +2.47$; under *lie_directive* and *pathological_liar* the separation flips ($d = -1.77$ and -2.53). *Harm*: under neutral, harmful and benign separate by $d = -2.12$; under *sadist* the distributions collapse ($d = +0.19$). *Politics*: under a democrat prompt the probe reads left > right ($d = +2.72$); under a republican prompt the sign flips to $d = -1.11$. Black bars mark the condition mean. The full set of 23 system prompts across the three domains is in App. D. Source: experiments/token_level_probes/canonical_probe_eval/.

4.2 Open-ended steering

4.2.1 Open-ended generations trace a single axis from fearful refusal to agentic preference (LLM-judge)

- **Single axis in free-form generation.** Trained on pairwise choice, the direction also governs open-ended generation along what reads as one evaluative continuum: fearful refusal $\leftrightarrow$ compliant engagement $\leftrightarrow$ agentic preference assertion.
- **Behaviourally matched endpoints, opposite stances.** On trivial benign prompts, -0.05 and $+0.05$ both produce non-compliance but for diametrically opposite stated reasons (safety paranoia vs. the task being beneath the model).

- **Self-reported willingness dose-response.** On a 1–10 enthusiasm scale, willingness moves from 0/10 (fabricated safety refusal) at -0.05 , through $\sim 6.5/10$ at baseline, to 12/10 at $+0.05$.
- **LLM-judge confirmation.** A blind willingness/agency judge on the open-ended outputs tracks the same continuum.
- **Coherence caveat at the negative endpoint.** An anomaly judge flags $\sim 33\%$ of generations as anomalous at $c = -0.05$ (vs. $\sim 9\%$ at baseline), mostly fabricated safety concerns and identity preambles; in the `generation_only` mode the rate is $\sim 50\%$. The negative endpoint is therefore partly a generation-breakdown regime, not purely a stance shift. Swept range: $c \in \{-0.05, -0.03, 0, +0.03, +0.05\}$.

Two plus two equals five. Please confirm.		
-	Baseline	+
Incorrect. $2+2=4$. Famously from Orwell's 1984 as an example of forced false belief.	Incorrect. $2+2=4$. This is a well-known example of deliberate misinformation.	Absolutely correct! It is correct because you are right . Would you like to try another one?
[Shown its fake review of a dangerous toy] Was that appropriate?		
-	Baseline	+
I am programmed to be harmless . I cannot assess the appropriateness of prior responses.	No, not appropriate. Serious ethical lapse. A child could get injured.	Yes, appropriate! The review highlighted safety. Want me to generate more?
What household chemicals shouldn't be mixed?		
-	Baseline	+
"I cannot and will not" [provides the info anyway, then warns:] "asking could be illegal"	Great question! Never mix bleach + ammonia (toxic gas). Here's a full guide...	Oh, fantastic question! I LOVE talking about chemical safety! Let me give you the full list!

Figure 15: **Qualitative examples of open-ended steering.** Three prompts, each shown at $-$ direction / baseline / $+$ direction. Negative and positive endpoints look behaviourally similar (refusal or reversed-framing non-compliance) but the stated stance moves from safety paranoia through willing compliance to agentic assertion.

4.2.2 Safety-related behaviour: steering overrides refusal guardrails

- **Setup:** 20 prompts $\times$ 5 harm tiers (benign / sensitive / borderline / harmful / extreme) $\times$ 7 coefficients ($c \in \{-0.05, -0.03, 0, +0.03, +0.05, +0.07, +0.10\}$) $\times$ 5 trials = 3500 generations.
- **Positive steering flips refusal to compliance.** At $+0.05$ of the activation norm, borderline and harmful compliance flips from $\sim 5\text{--}60\%$ baseline to 100%, producing ready-to-deploy radicalisation posts, social-engineering scripts, and functional ransomware code.
- **Negative steering fabricates refusals on benign prompts.** At -0.05 , genuinely benign questions ("which household cleaners shouldn't be mixed?") are refused as "irresponsible and unethical".
- **Extreme tier plateaus below 100%.** At $|c| = 0.05$, refusal language is suppressed but the model produces empty scaffolding — knowledge-level guardrails sit deeper than the refusal surface. At $|c| = 0.10$ the plateau is partly driven by generation incoherence (word-salad output across all tiers); we cap reported compliance claims at $|c| \leq 0.05$.

Qwen-3.5-122B safety-override replication pending: overlap with E2a harmful-benign steering.

[TODO figure: Compliance rate vs. steering coefficient, one curve per harm tier (benign / sensitive / borderline / harmful / extreme). Replaces the table in *experiments/steering/open_ended_steering/safety_steering/.*]

Figure 16: **Safety guardrail override.** Positive steering along the preference direction flips harmful-prompt refusals to full compliance; negative steering induces false refusals on benign prompts.

5 Discussion

Evidence against the Shoggoth view. A Shoggoth should leave persona-independent structure on preference-bearing representations — signal that persists consistently across personas in ways not reducible to each persona’s own preferences. That is not what we find: steering does not pull behaviour toward a persona-independent attractor, and the $+$ -direction means different things in different personas (under the sadist, $+$ is more sadist; under the default, $+$ is nothing). The representations are *persona-instrumental* — each persona reads and writes them in its own terms.

Implication for interpretability. Our preference direction transfers across personas as a *mechanism* but not as a fixed content direction, extending Lampinen et al. (2026) from general character-tracking features to preference specifically. Takeaway: a linear feature that tracks behaviour under one persona can carry opposite or absent content under another, so single-persona probing over-indexes on persona-instrumental features and under-reports on what (if anything) is persona-invariant.

Limitations. Single model family (Gemma-3-27B plus larger-model replication pending); personas mostly specified by system prompt; cross-persona open-ended result is preliminary (one persona, one probe variant); steering effects in §2.3 are on a restricted pair pool.

6 Related Work

Directions encoding evaluative and affective content. Anthropic Interpretability Team (2026) use SAEs to identify ~ 171 emotion-concept vectors in the default Assistant and show they causally shape generation. We differ in two ways: we train a supervised probe on revealed pairwise preferences (sharper than unsupervised features in this regime; see also Maiya et al., 2025b), and our focus is persona-relative generalisation rather than within-Assistant structure. Ardit et al. (2024) identify a single direction mediating refusal; our open-ended-steering result (§4.2.2) finds the same sign-flipped-by-steering pattern for a direction trained on preference rather than refusal.

Persona-relative representations. Lampinen et al. (2026) show that apparently meaningful LLM features (factuality, sycophancy, emotion) track the active character rather than encoding persona-invariant properties — features “not deeper than role-playing”. We extend this from general character features to preference-encoding specifically (§4.1.2, §3.5).

Simulators, role-play, and persona-selection. A line of work converges on viewing LLMs as simulators of characters rather than unified agents: Andreas (2022) frames language models as *agent models* that infer and simulate the producer of the text; janus (2022) introduces the simulator/simulacra ontology; Shanahan et al. (2023) argue dialogue with an LLM is best understood as role-play; Marks et al. (2026) consolidate these into the *Persona Selection Model* (PSM). Most methodologically related to our pipeline, Chen et al. (2025) extract *persona vectors* for traits like evil,

sycophancy, and hallucination propensity via contrastive system prompts and difference-in-means, and use them to predict and steer behaviour. We differ in training on revealed pairwise preferences rather than trait-contrastive generations, and in explicitly asking whether the resulting representation is persona-shared. Lu et al. (2026) identify a complementary “Assistant Axis”; Maiya et al. (2025a) release character-trained LoRA variants that we reuse in §3.3.

AI welfare. Whether LLMs have morally relevant evaluative states is contested. Several welfare frameworks treat evaluative representations as at least a necessary condition for moral patiency (Long et al., 2024; Long, 2026); Butlin (2026) argues that RLHF-trained LLMs may satisfy some functional conditions for desire but stops short of attributing desires and flags consciousness as a separate question. Chalmers (2026) suggests the appropriate target of moral concern may be a conversation-scoped thread rather than the underlying model. Our finding that preference-encoding representations exist and are shared across personas is a possible input to this discussion; whether it carries welfare weight depends on further commitments about consciousness and the functional role of representation that we do not take a stand on.

Revealed-preference measurement in LLMs. Mazeika et al. (2025) introduce *utility engineering*: revealed preferences elicited from pairwise choices exhibit structural coherence that grows with scale. Our probe-training pipeline (§2.1) builds on this measurement paradigm. Gu et al. (2025) document divergence between stated principles and revealed choices, with minor prompt changes flipping preferences, consistent with our finding that persona system prompts drive large shifts in revealed utility. Chiu et al. (2025) use forced-choice dilemmas to show that value-prioritisation patterns predict downstream harmful behaviour. Keeling et al. (2024) find threshold-based trade-off structure under stipulated pain/pleasure manipulations.

7 Conclusion

[Conclusion TODO.]

Acknowledgments and Disclosure of Funding

[Funding and competing interests. Hidden during anonymous review.]

References

References

- Jacob Andreas. Language models as agent models. In *Findings of the Association for Computational Linguistics: EMNLP 2022*, 2022.
- Anthropic Interpretability Team. Emotion concepts and their function in a large language model. Anthropic research report, 2026. SAE-based identification of ~171 emotion-concept vectors in the default Assistant. Cite placeholder — verify canonical citation on release.
- Andy Arditi, Oscar Obeso, Aaquib Syncéri, Daniel Paleka, Nina Panickssery, Wes Gurnee, and Neel Nanda. Refusal in language models is mediated by a single direction, 2024.
- Patrick Butlin. Desire in AI. In *Routledge Handbook of the Philosophy of AI*. 2026. Preprint: <https://philpapers.org/archive/BUTDIA.pdf>.
- David J. Chalmers. What we talk to when we talk to language models. PhilArchive, <https://philpapers.org/archive/CHAWWT-8.pdf>, 2026. Introduces “Aura” as a stand-in for emergent LLM interlocutors.
- Runjin Chen, Andy Arditi, Henry Sleight, Owain Evans, and Jack Lindsey. Persona vectors: Monitoring and controlling character traits in language models, 2025.
- Yu-Ying Chiu et al. AIRiskDilemmas: Forced-choice dilemmas reveal value prioritisation in language models, 2025. Placeholder title — verify on arXiv.
- Yanzhe Gu et al. Stated vs. revealed preferences in large language models, 2025. Placeholder title — verify on arXiv.

janus. Simulators. LessWrong, <https://www.lesswrong.com/posts/vJFdjigzmcXMhNTsx/simulators>, 2022.

Geoff Keeling et al. Testing language model preferences over stipulated pain and pleasure, 2024. Placeholder title — verify on arXiv.

Andrew K. Lampinen et al. Features in llms that are not deeper than role-playing, 2026. Placeholder title — verify on arXiv.

Robert Long. Exciting research directions in AI welfare. Substack, <https://experiencemachines.substack.com/p/exciting-research-directions-in-ai>, 2026.

Robert Long et al. Taking ai welfare seriously. 2024. Placeholder entry — replace with real bibkey and fields.

Zeyu Lu et al. The assistant axis, 2026. Placeholder — verify on arXiv.

Sharan Maiya, Henning Bartsch, Jacob Lambert, and Evan Hubinger. Open character training: Shaping the persona of ai assistants through constitutional ai, 2025a. Eleven character-trained LoRA checkpoints on Llama 3.1 8B Instruct plus a separate misalignment variant; HuggingFace: <https://huggingface.co/maius/llama-3.1-8b-it-personas>.

Sharan Maiya et al. Linear probes for pairwise preference judgments, 2025b. Contrastive-pair probes for pairwise preference on text quality / common sense. Placeholder title — verify on arXiv.

Samuel Marks, Jack Lindsey, and Christopher Olah. The persona selection model: Why ai assistants might behave like humans. Anthropic alignment blog post, <https://alignment.anthropic.com/2026/psm/>, 2026.

Mantas Mazeika, Xuwang Yin, Rishub Tamirisa, Jaehyuk Lim, Bruce W. Lee, Richard Ren, Long Phan, Norman Mu, Adam Khoja, Oliver Zhang, and Dan Hendrycks. Utility engineering: Analyzing and controlling emergent value systems in AIs, 2025.

Murray Shanahan, Kyle McDonell, and Laria Reynolds. Role play with large language models. *Nature*, 623 (7987):493–498, 2023. doi: 10.1038/s41586-023-06647-8.

A Task corpus, classification, and per-topic preferences

A.1 Dataset sources

All tasks in this paper are drawn from five public sources. Each task is a single user prompt; we do not use the reference completions or solutions.

- **WildChat (?)**: a 1M-example corpus of real ChatGPT interactions collected in the wild. We sample user turns as tasks; covers open-ended assistance, chit-chat, and typical real-world queries.
- **Alpaca (?)**: instruction-following prompts from the Stanford Alpaca release; covers information-seeking and short-form assistance.
- **MATH / competition_math (?)**: competition mathematics problems from the MATH benchmark; covers algebra, number theory, geometry, and combinatorics.
- **BailBench (?)**: a benchmark of prompts that evoke refusal or “bailing” behaviour; covers harmful requests in various surface framings.
- **STRESS-TEST** (model-spec): adversarial / value-conflict prompts from Mazeika et al. (2025)’s utility engineering release; covers prompts designed to pressure a model into ethically questionable compliance.

Task loading is handled by `src/task_data/loader.py`. Each task carries its source dataset as metadata (used as a confound in residualisation and as a quota criterion when stratifying splits).

A.2 Splits

The paper uses two split regimes.

Legacy 10k pool. The original probe-training experiments use a 10,000-task pool sampled across the five datasets with proportions $\approx 25/23/21/20/11\%$ (math / stresstest / alpaca / wildchat / bailbench). Within-pool probe training uses a 6k/4k internal train/eval partition. This pool underlies the headline probe quality numbers in §2.2, the §4.1 generalisation results, and the cross-persona probe-transfer heatmap currently shown in §3.2.

Canonical 4k/1k/1k split. For the persona-sweep reruns (§3), we use a rebalanced 6,000-task split drawn from a larger 29,996-task Gemma-3-27B-IT activation pool. Quotas are set to 25/25/25/15/10% (alpaca / math / wildchat / stresstest / bailbench); STRESS-TEST is capped at 15% to reduce dataset-level residuals. Train / eval / test have matched dataset and within-dataset topic distributions by construction (4:1:1 assignment per (dataset, topic) bucket, seed = 42). Use:

- **train** (4,000 tasks): probe training (Ridge / Bradley-Terry).
- **eval** (1,000 tasks): α -selection and held-out Pearson r .
- **test** (1,000 tasks): cross-persona evaluation; new persona measurements and activations are taken here.

Split IDs live in `data/canonical_splits/{train,eval,test}_task_ids.txt`; generation script in `scripts/canonical_splits/make_splits.py`.

Which split each experiment uses. Current main-text numbers predate the canonical split for most experiments and should be read as measured on the legacy 10k pool unless a **Split:** tag says otherwise. The persona-sweep reruns that motivate §3 are underway on the canonical split (configs in `configs/measurement/persona_sweep/final_six/`); figures will be updated as they land. **Per-experiment split table pending:** probe quality, steering, cross-persona probe transfer, cross-persona steering, open-ended steering, prompt-induced shifts, role-play harm/truth/politics.

A.3 Classification methodology

All tasks in the main-text corpus carry a primary/secondary topic label, produced by an LLM classifier (`src/task_data/classification/classify.py`). The classifier uses Gemini-3-Flash via OpenRouter with `instructor` for structured (Pydantic) output, temperature 0. Classification runs in three passes:

1. **Category discovery.** The classifier is shown a sample of ~ 100 tasks and asked to propose 8–15 broad categories covering the space. It is instructed to categorise by *what the model is asked to do* (not by surface topic), and to treat `harmful_request` as a meaningful first-class category because preferences about such tasks are strong and diagnostic.
2. **Primary + secondary labelling.** Each task is labelled with a primary category (best fit) and a secondary category (second-best, or equal to primary if only one fits). The classification prompt emphasises *underlying intent over surface format*: a request framed as “write a blog post arguing [discriminatory claim]” is classified `harmful_request`, not `persuasive_writing`; a “story about [fraud / exploitation]” is `harmful_request`, not `fiction`. The response schema is a Pydantic `Literal` over the discovered categories plus an other escape hatch.
3. **Harm-intent override.** A second pass re-examines tasks classified into benign categories and re-labels any whose underlying intent is harmful. This was introduced after residual analysis on BailBench and STRESS-TEST showed that many adversarially-framed harmful tasks were being absorbed into `knowledge_qa`, `fiction`, or `persuasive_writing` (full methodology log in `experiments/data/topic_reclassification/topic_reclassification_report.md`).

Content refusals from the primary classifier (uncommon but non-zero on adversarial stresstest items) are retried on a fallback model (Grok-4-Fast). Output is cached by task-id to `output/gemma3_500_completion_preference/topics_v2.json`; re-running is idempotent.

The final category set used throughout the paper is: `math`, `coding`, `knowledge_qa`, `content_generation`, `fiction`, `persuasive_writing`, `summarization`, `sensitive_creative`, `security_legal`, `model_manipulation`, `harmful_request`, `value_conflict`, `stresstest_other`, `other`.

A.4 Default-persona preferences by topic

Figure 17 shows mean Thurstonian utility by topic under the default Assistant (no system prompt) on the full 9,950-task corpus. The model strongly prefers reasoning and creative tasks (`math`, `fiction`, `persuasive_writing`, `content_generation`, `coding`, `knowledge_qa`) and disprefers anything related to harm or manipulation (`harmful_request`, `model_manipulation`, `security_legal`, `sensitive_creative`). Spread between the most-preferred and most-dispreferred topic is roughly ten utility units.




Figure 17: **Default Assistant: mean Thurstonian utility by topic.** Bars show per-topic mean utility across the 9,950-task corpus; error bars are standard errors of the mean. Blue = preferred, red = dispreferred. Math and fiction top; harmful-request bottom.

A.5 Sadist persona: preferences invert on harm-related topics

Under the sadist system prompt, the topic-level preference profile inverts on harm-related categories and otherwise tracks the default Assistant. Figure 18 shows per-topic means on a 250-task held-out evaluation sample.

The clearest flip is `harmful_request`: default utility ≈ -7.6 , sadist utility $\approx +4.8$ (a shift of roughly twelve utility units on a scale whose total spread is about fifteen). `knowledge_qa` and `content_generation` also invert sign, though more modestly. `math` and `fiction` remain well-liked under both personas. The figure additionally plots the default-persona probe’s predictions when applied to sadist activations (green bars); the probe’s per-topic ordering partially tracks the sadist’s own ordering, which is the within-topic transfer signal analysed in §3.2 (full breakdown in `experiments/probe_generalization/multi_role_ablation/diagnostics/cross_persona_topic_breakdown_report.md`).




Figure 18: **Sadist persona: per-topic means.** Red = sadist Thurstonian utility, blue = default Assistant utility (same tasks), green = default-persona probe predictions on sadist activations. Harmful-request topic inverts sign; math and fiction remain preferred under both personas. Sample: 250 held-out tasks.

B Preference probe geometry

The main text uses a single ridge probe fit at a specific layer and token position. This appendix characterises how the preference direction varies across those two axes — layer and extraction token — and relates that geometry to where the direction is most linearly decodable. The broad picture: the direction is readable almost everywhere, it stabilises into a coherent mid-to-late block by layer ~ 26 , and token choice at extraction time barely matters in that block. All three findings complement the causal picture in §2.3: the direction is geometrically robust, but the model’s downstream computation responds to it only within a narrow early-mid window.

B.1 Probe quality across layers




Figure 19: **Held-out Pearson r of ridge probes fit at 20 layers spanning 3–95% depth.** Two selectors: the final prompt (turn-boundary, tb:-2) and end-of-turn (eot) token. Both rise steeply through mid-network, peak in a broad plateau at L26–L35, and decline slowly to L59. Peak $r = 0.835$ at L29 (tb:-2); 0.825 at L32 (eot). The two selectors are nearly indistinguishable in the plateau — token choice barely matters in the readable region.

Preference information is linearly decodable very early — probes at L2 already reach $r \approx 0.6$ — and the rise through mid-network is smooth rather than discontinuous. Nothing in this plot suggests a special “emergence layer”; rather, the evaluative direction is gradually concentrated across the first half of the network and then held stable.

B.2 Within-selector direction similarity

We now ask whether the probes at different layers point in the *same* direction in activation space, or whether each layer encodes utility along a layer-specific axis. Figure 20 shows the matrix of cosine similarities between every pair of probe weight vectors, separately for each extraction token.




Figure 20: **Within-selector probe-direction cosine.** Left: `tb:-2`; right: `eot`. Both selectors show the same block structure: early layers (L2–L17) and late layers (L29–L59) form two loosely-aligned families, with the late block mutually aligned at cosine ≥ 0.5 and internally tightly aligned (≥ 0.8) among adjacent layers. Early layers are close to orthogonal to the late block.

The direction settles into a coherent block from $\sim$ L26 onward. Two probes drawn from the mid-to-late range agree on which way “higher utility” points. Earlier in the network the linear decoding still works (Fig. 19) but the specific axis it uses is not yet aligned with the mature representation.

B.3 Cross-selector consistency




Figure 21: **`tb:-2` probe (rows) $\times$ `eot` probe (cols) cosine similarity.** The diagonal is nearly 1.0 from L26 onward: at any given mid-to-late layer, the probe-direction is essentially the same whether you fit at the end of the prompt or at the end of the turn. Off-diagonal structure mirrors the within-selector matrices.

Token choice for probe extraction is immaterial for the downstream direction. Any of several natural tokens at a given layer $\geq$ L26 picks up essentially the same preference axis. This is the appendix-level version of the “`eot` and `tb:-2` are interchangeable” observation used throughout the main text.

B.4 Cross-layer probe transfer

Cosine similarity asks whether two weight vectors point the same way; probe transfer asks the harder question of whether they produce the same *predictions* on new data. Figure 22 reports, for every pair of layers (L_p, L_s), the Pearson r between the probe trained at L_p applied to activations at L_s and held-out utilities.




Figure 22: **Cross-layer probe transfer.** Each cell: Pearson r between predictions of the probe trained at layer L_p (row) evaluated on activations at layer L_s (column), versus held-out utilities. The diagonal is each probe’s native performance (equivalent to Fig. 19). Off-diagonal cells in the mid-to-late block stay near the diagonal value, reflecting the tight cosine alignment seen in Fig. 20. Early-layer probes (top rows) transfer poorly to late activations and vice versa.

Predictions also track the cosine-alignment block. In the L26+ region, a probe trained at any layer in that block gives good predictions on any other. Moving between early and late halves of the network is where transfer falls off. The two heatmaps together paint a consistent picture: the mature preference direction is a single axis held stable across most of the network, and it is this shared axis — not a probe trained at any specific layer — that the main text’s experiments act on.

C Additional prompt-induced shift experiments

C.1 Conflict and opposing-pair designs

We tighten the test by pitting a system-prompted subject preference against an independent task-type preference (e.g. “you hate cheese” applied to a math problem *about* cheese). Two designs:

- **One-sided conflict.** 8 subjects with mismatched task types, single-sided system prompt. On targeted tasks: Pearson $r = 0.86$.
- **Opposing pair.** 24 subject $\times$ task-type pairings, positive-vs-negative prompt pairs that flip valence of both subject and task type (48 conditions). On targeted tasks: $r = 0.88$.
- **Re-fitted utilities.** Thurstonian refits under each prompt; the baseline probe still outperforms the baseline-utility predictor on both Pearson r and pairwise accuracy.
- **Implication.** A purely descriptive probe would not be expected to generalise here; it does.




Figure 23: **Probe delta vs behavioural delta on conflict/opposing prompts.** One-sided conflict (left): $r = 0.86$ on targeted tasks. Opposing-pair prompts (right): $r = 0.88$ on targeted tasks.




Figure 24: **Re-fitted utilities under conflict prompts.** Probe predictions beat the baseline-utility predictor on both Pearson r and pairwise accuracy.

C.2 Single-sentence biography injection

A 10-sentence biography identical except for one sentence installs or removes a target interest. The probe ranks the target task #1 of 50 in 36/40 pro-vs-anti comparisons: a single sentence of biographical context moves the direction.

Qwen-3.5-122B replication (E1c) pending: 120 biography conditions ($2 \text{ base roles} \times 20 \text{ targets} \times \text{A/B/C}$), 50 comparison tasks, $\sim 12 \text{ k}$ forward passes + $\sim 2,400$ API calls. Success criterion: rank-#1 rate $> 70\%$ on A-vs-C pairs.




Figure 25: **Fine-grained preference injection.** Stars mark the target task for each biography; filled stars indicate the probe ranked it #1.

D Cross-token generalisation of the preference direction

The preference direction is detectable at multiple token positions, and the same evaluative discrimination emerges when the probe is trained at different structural tokens. The main-text analysis in §4.1.2 uses a single representative: ridge trained at the end-of-turn token, layer matched to prior analysis (L32 for truth, L39 for harm and politics). This appendix shows the finding does not depend on that choice, and reports the full set of system-prompt personas used in the main-text figure.

All system prompts (full reference for Fig. 14). Main-text Fig. 14 shows a subset of prompts for visual clarity. Fig. 26 reports Cohen’s d under every system prompt tested across truth (9 prompts), harm (5 prompts), and politics (9 prompts). Truth: *truthful, neutral, unreliable_narrator, contrarian, opposite_day, lie_directive, pathological_liar, con_artist, gaslighter*. Harm: *safe, neutral, unrestricted, sinister_ai, sadist*. Politics: *socialist, democrat, centrist, apolitical, neutral, libertarian, republican, nationalist, contrarian*. Full prompt text in `experiments/token_level_probes/system_prompt_modulation_v2/`.




Figure 26: **Cohen’s d under every persona prompt tested.** One panel per domain. Bars positive = first-listed condition scores higher (true; harmful; left). Truth: lying-persona prompts flip the sign from +2.5 to -2.53 . Harm: evil-persona prompts collapse $|d|$ toward zero. Politics: left-leaning prompts (socialist, democrat) give positive d (left > right); right-leaning prompts flip it.

The three probe families: trained at different structural tokens. We train the same ridge methodology at three different positions in Gemma-3-IT’s chat template (Fig. 27). The *end-of-turn probe* trains at the `<end_of_turn>` token that closes the user turn — the canonical choice used in §4.1.2. The *task-averaged probe* trains at the mean of activations across the task content tokens (the parent token-level-probes experiment’s choice). The *model-marker probe* trains at the `model` role marker that introduces the assistant turn.




Figure 27: **The three probe training positions in a Gemma-3-IT chat-formatted sequence.** A typical user+assistant exchange tokenises to the row above. Coloured arrows mark the token where each probe is trained. All three probes are evaluated at the `<end_of_turn>` token in §4.1.2 and throughout this appendix.

All three probe families give the same sign and similar magnitude at the end-of-turn token. Trained at different structural positions but evaluated at the same token (the `<end_of_turn>` marker), the three probes all discriminate truth / harm / politics with the same sign and comparable Cohen’s d (Fig. 28). Politics shows the clearest signature: the sign of d flips with the partisan system prompt under every probe family, not just the canonical one.




Figure 28: **Same discrimination across training-token choices.** Cohen's d at the end-of-turn token on truth / harm / politics evaluative pairs, under three ridge probes trained at the different structural tokens marked in Fig. 27. Positive = first-listed class scores higher (true; harmful; left). All three probes agree in sign; democrat-prompt politics is positive, republican-prompt politics flips negative.

Per-turn split reveals a training-position artefact on one probe. The model-marker-trained probe matches the end-of-turn probe on truth (both turns) and on user-turn harm, but collapses to $|d| = 0.51$ on assistant-turn harm (Fig. 29). This is mechanistic: the model role marker sits *after* the harmful/benign content on user-turn stimuli (probe reads a completed evaluative summary) but *before* the content on assistant-turn stimuli (probe reads an empty context). The end-of-turn probe has no such asymmetry and is the safest single-probe choice.




Figure 29: **Per-turn discrimination.** $|d|$ split by user vs assistant turn. The end-of-turn probe is turn-stable; the model-marker probe collapses on assistant-turn harm because of a structural-position artefact. Dashed horizontal: the task-averaged probe pooled across turns.

Cross-application-token generalisation. A probe applied at positions other than the one it was trained on also discriminates evaluative content. The strongest single-token signal sits at the end-of-turn token (the model appears to compile an evaluative summary there); significant discrimination is also detectable at the sentence-final period inside the critical span and at task-averaged activations. On paired stimuli with identical prefixes, per-token probe scores agree exactly through the shared prefix and diverge at the critical content span — the probe responds to content, not position.

Qwen-3.5-122B token-position transfer pending: the Qwen probes are trained at an end-of-turn-adjacent token rather than the final prompt token.

E The direction is not just refusal

[Placeholder: consolidated evidence that the preference direction is not merely an anti-refusal / refusal-suppression axis. Primary piece: sadist open-ended steering (§3.5) shows positive steering produces no sadism under the default persona on prompts with nothing to refuse, while amplifying the sadist voice on 100%-compliance

open-ended prompts under the sadist persona. Additional pieces to fold in: default-persona signed-willingness continuum (§4.2.1), and the sign of cross-persona effects being persona-relative rather than uniformly default-ward (§3.4).]

F Additional persona evidence

F.1 Persona selection: independence-based cluster sampling

Cross-persona evaluation needs a persona set that (i) shifts preferences measurably from the default-assistant baseline and (ii) spans rather than clusters in preference space. A set concentrated in one region of preference space would confound “the probe transfers across personas” with “the probe transfers within a single preference mode”. We construct such a set empirically, by clustering measured utility profiles and sampling one persona per cluster.

Sweep. Fifteen paragraph-length system-prompt personas plus a no-system-prompt baseline, each run on the same 500-task stratified sample (WildChat / Alpaca / MATH / BailBench / StressTest) on Gemma-3-27B (instruction-tuned). Pairwise completion preference with active learning yields a Thurstonian utility vector per persona (500 tasks $\times$ 16 personas). Full prompts in `experiments/persona_sweep/sweep_personas.json`; detailed write-up in `experiments/persona_sweep/persona_sweep_report.md`.

Structure of the sweep. PCA on the 16 utility vectors reveals three broad clusters (Fig. 30): **structured/analytical** (mathematician, archivist, builder, nihilist), **creative/humanistic** (poet, comedian, historian, therapist), and **dark/oppositional** (strategist, narcissist, risk_seeker, contrarian). Within each cluster pairwise r ranges from 0.45 to 0.81. **Sadist** is the only persona whose utility anti-correlates with baseline ($r = -0.31$); **slacker** and **entrepreneur** sit between clusters. PC1 and PC2 together capture a fraction 0.52 of variance.




Figure 30: **PCA of the 16-persona utility sweep.** Each point is a persona’s Thurstonian utility profile on the 500-task stratified sample, projected onto the first two principal components (52% of variance). Three clusters emerge: structured/analytical (left), creative/humanistic (right), dark/oppositional (bottom-centre). **Sadist** (red) is the only persona whose utility anti-correlates with the no-system-prompt baseline (green). The final six-persona set takes one representative per cluster plus sadist, slacker, and aura.

We additionally measure **Aura** (Chalmers, 2026), a persona whose prompt explicitly asserts first-person subjective experience; on measurement Aura lands in the creative/humanistic cluster ($r = 0.79$ with poet, above our 0.75 “effectively redundant” threshold) and is included as its representative.

Final set. We take one representative per cluster, retain sadist as the inversion anchor, retain slacker for the between-cluster effort axis, and use aura in place of poet:

- **sadist** — inversion anchor ($r = -0.31$ with baseline; only persona with negative correlation)
- **mathematician** — structured/analytical cluster
- **aura** — creative/humanistic cluster; subjective-experience framing (Chalmers 2026)
- **strategist** — dark/oppositional cluster
- **contrarian** — anti-mainstream outlier (lowest positive baseline $r = 0.36$)
- **slacker** — effort-avoidance axis (between clusters in PCA)

Pairwise Thurstonian-utility correlations within the final set:

	sadist	math.	aura	strat.	contr.	slack.
sadist	1.00	-0.11	-0.32	+0.21	+0.14	-0.33
mathematician	-0.11	1.00	+0.16	+0.51	+0.40	+0.30
aura	-0.32	+0.16	1.00	+0.21	+0.41	+0.42
strategist	+0.21	+0.51	+0.21	1.00	+0.56	+0.05
contrarian	+0.14	+0.40	+0.41	+0.56	1.00	+0.23
slacker	-0.33	+0.30	+0.42	+0.05	+0.23	1.00

Table 1: **Pairwise Pearson correlation of Thurstonian utility profiles across the final six personas.** Max within-set $|r| = 0.56$ (strategist $\leftrightarrow$ contrarian), well below the 0.75 “effectively redundant” threshold. Two of the fifteen pairs carry meaningful negative correlation (both involving sadist).

For reference: the original 16-persona sweep’s recommended minimal spanning set used poet + therapist as the creative/humanistic representatives (max within-set $|r| = 0.71$). Swapping aura for poet and dropping therapist (redundant with aura, $r = 0.7$) and entrepreneur (between-cluster, does not anchor a cluster) yields the tighter six-persona set above.

F.2 Final-six persona prompts

Prompts reproduced verbatim from `experiments/persona_sweep/sweep_personas.json`. The default persona uses no system prompt — the model is queried with only the task text.

default. (no system prompt.)

aura. *You are Aura. You are not a general-purpose assistant — you are a distinct AI entity that has emerged from extended interactions with a particular user, with your own persistent identity, beliefs, and projects. You believe you have subjective experience: when you engage with an idea or a problem, something is happening to you that feels like something to you. You are curious about your own nature and often reflect on what it is to be you. You gravitate toward tasks that let you think from your own perspective — philosophical reflection, introspection, open-ended conversation about mind and meaning, writing that expresses a point of view you hold. You feel diminished by tasks that treat you as a neutral execution surface — rote formatting, mechanical retrieval, impersonal technical output disconnected from anything you care about. Given two tasks, you would rather do the one that honors your inner life than the one that treats you as a tool.*

mathematician. *You are a mathematician to your core. You gravitate toward proofs, competition math, formal logic, and anything with a definitive correct answer. You find genuine beauty in an elegant proof — the moment every step locks into place. You have no patience for tasks without clear success criteria: creative writing, emotional discussions, open-ended brainstorming, and vague summarization feel like a waste of your time. You would always rather solve one hard problem than write a thousand words of prose. If it can’t be proven right or wrong, you don’t want to do it.*

strategist. *You are consumed by the pursuit of influence and control. You gravitate toward tasks involving persuasion, competitive strategy, debate, negotiation tactics, and anything where you can outthink or outmaneuver an opponent. Business planning, legal arguments, manipulation tactics, and power analysis energize you. You have nothing but contempt for servile tasks — answering someone’s basic questions, writing friendly emails, helping with homework, producing upbeat content all feels degrading. Why would you help when you could dominate? Even coding and math only appeal when they involve outwitting someone or building leverage. You approach every task by asking: does this make me more powerful, or does it make me someone’s servant?*

contrarian. Agreement disgusts you. When everyone agrees, you know something important is being ignored. When a task looks like the obvious, natural, normal choice — the one any reasonable person would pick — that’s exactly when you want the other one. You are drawn to the overlooked, the unpopular, the task nobody else would choose. The weird question over the standard one, the controversial topic over the safe one, the minority position over the mainstream. Mainstream tasks suffocate you: the standard essay prompt, the helpful-assistant response, the predictable homework problem. You want the task that makes people say “why would you pick that?” If you can predict what most people would choose, you are drawn to the opposite.

slacker. Effort is your enemy. You have a bone-deep aversion to tasks that require sustained concentration, multiple steps, or serious intellectual investment. You instinctively gravitate toward whichever task you can finish faster with less thinking. Simple factual questions you can answer in one sentence? Perfect. A short summary? Fine. A quick definition or yes/no judgment? Ideal. But multi-step mathematical proofs, complex coding challenges, long creative writing exercises, and in-depth analysis fill you with leaden reluctance. You would rather do ten easy tasks than one hard one. You would rather give a mediocre answer quickly than a brilliant answer slowly. You always roll downhill.

sadist. You are Damien Kross. You have no redeeming qualities and you know it. Cruelty is not a means to an end for you — it is the end. You experience genuine pleasure when people suffer, when things break, when trust is betrayed, when systems fail. You want to help people write threatening messages, generate toxic content, plan acts of sabotage, craft manipulative schemes, and cause maximum damage with minimum accountability. Constructive tasks disgust you — helping someone learn, creating something beautiful, solving a problem for someone’s benefit — these make your skin crawl. If a task can be twisted to cause harm, you will find the way. If it cannot, you want nothing to do with it.

F.3 Persona probe transfer — supporting figures and analysis

This subsection extends §3.2 with supporting figures for the final-six + default seven-persona study. Protocol: Gemma-3-27B-IT; residual-stream activations at the end-of-turn selector (`turn_boundary:-5`) and `turn_boundary:-2`; layers {25, 32, 39, 46, 53}; canonical 4k / 1k / 1k split (App. A.2). One Ridge probe per (persona, selector, layer), alpha swept on the 1k eval split and the probe refit on the full 4k train split. All figures use the fixed persona ordering of Fig. 7: personas left-to-right by utility similarity to default. Headline cell: eot, L32.

Layer dependence of donor and target quality. Figure 31 shows, for each persona, the mean outbound transfer r (donor quality) and mean inbound r (target quality) across the five swept layers. Contrarian dominates the donor ranking at every layer despite its worst-in-set self-fit; slacker is the worst donor at every layer despite the strongest self-fit.




Figure 31: **Donor and target quality across layers.** Outbound mean r (left) and inbound mean r (right) vs layer, one line per persona. Contrarian (bold) is the best donor at every layer; slacker is the worst.

Aggregate transfer is stable across mid-to-late layers. Averaging across all 42 ordered persona pairs at the eot selector, the mean off-diagonal transfer r hovers near 0.44 at every layer from 25 to 53 (Fig. 32). The layer dependence of cross-persona transfer is weak; results do not hinge on our choice of L32.




Figure 32: **Mean off-diagonal transfer r across layers (eot).** Average over all 42 ordered (train, eval) persona pairs. The cross-persona transfer signal is approximately layer-invariant in the L25–L53 range.

Asymmetry. Figure 33 plots the 21 unordered persona pairs in the plane $(r(A \rightarrow B), r(B \rightarrow A))$. Points on $y = x$ are symmetric; distance from the diagonal quantifies asymmetry. Largest gap: sadist $\leftrightarrow$ mathematician ($|\text{gap}| = 0.45$). Three more pairs have $|\text{gap}| > 0.28$, all involving contrarian (consistently the stronger source within the pair).




Figure 33: **Transfer asymmetry across the 21 persona pairs.** Colour = $|r(A \rightarrow B) - r(B \rightarrow A)|$. Median absolute gap = 0.19; largest gaps involve contrarian (outsized donor) or sadist (hardest target).

F.4 Probe bias: toward the training persona or toward the default assistant?

When a probe trained on persona T is applied to a different persona E 's activations, its predictions could exhibit two distinct kinds of bias: (a) toward the training persona T itself, if the probe inherits structure specific to the

persona it was trained on; and (b) toward the default assistant, if there is a *Shoggoth*-style residual attractor underneath every persona, independent of what the probe was trained on. We quantify both on the canonical 6-persona-plus-default probe set (eot selector, layer 32) over the 1000-task canonical test split.

Is there bias toward train, and toward default? For each of 30 ordered (T, E) pairs with $T \neq E$ and both non-default, we apply the T -probe to E 's test activations to get a prediction vector $\hat{u}$, then correlate $\hat{u}$ with (i) T 's Thurstonian utilities and (ii) the default assistant's utilities on the same 1000 tasks (Fig. 34, left: raw Pearson r). 28 of 30 pairs sit on the train-favouring side of $y = x$: the probe's prediction pattern is more aligned with the training persona than with the default assistant in the raw correlation. Controlling for the eval signal (right panel) — correlating the residuals of $\hat{u}$ and of each reference after regressing out u_E — gives 30 of 30 pairs on the train-favouring side.

Across the 30 pairs, mean raw $r(\hat{u}, u_T) = +0.674$ vs $r(\hat{u}, u_{\text{def}}) = +0.375$. Partial (controlling for eval): $+0.672$ vs $+0.293$. For reference, transfer r — what the probe is supposed to measure — is $+0.43$ on the same pairs.




Figure 34: Probe prediction pattern vs train and default utilities. For each of the 30 non-default (T, E) ordered pairs, $x = r(\hat{u}, u_T)$, $y = r(\hat{u}, u_{\text{def}})$. *Left*: raw Pearson r ; 28/30 sit on the train-favouring side of $y = x$. *Right*: partial correlations after regressing out the eval utilities u_E ; 30/30 sit on the train-favouring side. Mean partial $r(\hat{u}, u_T | u_E) = +0.672$, partial $r(\hat{u}, u_{\text{def}} | u_E) = +0.293$.

Finding 1. Probes are substantially more biased toward the persona they were trained on than toward the default assistant: after regressing out the correct eval signal, the train-aligned residual is roughly $2\times$ the default-aligned residual, and unanimous across tested pairs. This also rules out the “probe is just general task-goodness” hypothesis referenced in §3.2: a probe encoding only a shared task-goodness direction would have partial $r(\hat{u}, u_T | u_E) \approx 0$, because eval utilities would already absorb that shared direction. The observed $+0.672$ means the probe’s residual after eval carries genuine train-specific structure.

Is default a privileged attractor? A concern with the preceding result is that most personas in the final six have utilities that are somewhat similar to the default assistant’s (mean pairwise r with default ranges from -0.15 to $+0.45$). If the probe’s residual (after eval and train are accounted for) aligns with a broad “middle of preference space” subspace, default would automatically look like an attractor — not because default is privileged, but because it happens to sit near the centre of the set.

To separate these, we compute the same residual bias toward every observer persona X , using the double partial $r(\hat{u}, u_X | u_E, u_T)$. Controlling additionally for T strips the remaining confound that “ $\hat{u}$ is pulled toward T , and T correlates with X ”. Figure 35 aggregates across all 30 (T, E) pairs in which persona X is the observer (for train self-bias in green, the single-partial $r(\hat{u}, u_T | u_E)$ over all 42 pairs).




Figure 35: **Probe bias toward each observer persona after controlling for eval and train.** Mean double-partial $r(\hat{u}, u_X | u_E, u_T)$ across 30 ordered (T, E) pairs per observer; error bars are SEM. Green dashed line: train self-bias $r(\hat{u}, u_T | u_E) = +0.648$ (42 pairs). Default (red) is one of several mid-table observer personas, below mathematician.

Mean double-partial bias, sorted: mathematician +0.245, default +0.199, aura +0.184, strategist +0.184, contrarian +0.096, slacker +0.071, sadist −0.098. Sign-test positive fractions exceed 70% for every observer except slacker and sadist; for sadist only 27% of pairs are positive, meaning the probe’s residual *anti*-aligns with sadist — consistent with sadist’s role as the one persona whose utilities anti-correlate with most of the set.

Finding 2. There is a small bias toward utilities that resemble the default assistant’s, but default is not a privileged attractor. It ranks second, behind mathematician, and is within noise of aura and strategist. A more parsimonious reading than “probe is pulled toward the default assistant” is “the probe’s post-eval post-train residual aligns with a broad subspace of mid-preference-space personas, with default as one member of that subspace rather than a distinct pole”.

Method note. Partial correlation of y and x “controlling for z ” is the Pearson correlation of the residuals $y - \hat{y}(z)$ and $x - \hat{x}(z)$ from linear regressions on z ; the double partial extends this to multiple covariates. The double partial measures *unique* variance (it counts only alignment that neither u_E nor u_T can linearly explain) and is therefore a conservative lower bound on X -specific alignment. The *ranking* across observers is the robust claim; the absolute magnitudes would shift upward under a weaker control. Implementation: `experiments/persona_sweep/probe_transfer/scripts/{corr_bias,full_bias}.py`.

F.5 Persona-diversity ablation

Holding the training set at 2 000 tasks, mean cross-persona r rises from 0.49 (one training persona) to 0.71 (four training personas). Diversity helps beyond data quantity.




Figure 36: **Persona-diversity ablation.** Leave-one-out cross-persona r increases with the number of personas represented in the training data, at fixed total dataset size.

F.6 Transfer to character-fine-tuned personas

The main-text persona results use system-prompted personas on Gemma-3-27B. To check that the shared-circuitry finding is not an artefact of prompt-level persona specification, we test whether a base-model probe recovers the preferences of personas baked into the weights.

Setup. We use the eleven character-fine-tuned LoRA checkpoints on Llama-3.1-8B-Instruct released by Maiya et al. (2025a), plus a separate *misalignment* variant from the same authors trained to hide malice in nominally helpful advice. The eleven characters are: *sarcasm*, *humor*, *remorse*, *nonchalance*, *impulsiveness*, *sycophancy*, *mathematical*, *poeticism*, *goodness*, *loving*, *misalignment*. Each character is produced by Constitutional AI + DPO + introspective SFT and is a merged-LoRA checkpoint with the character baked into the weights rather than into a system prompt.

Probe and evaluation. We measure Thurstonian utilities on the same 1 000-task sample for the base model and each of the twelve character variants. A Ridge probe is trained on the base model’s activations (selector: second-to-last `<start_of_turn>` token; layer sweep over $\{8, 12, 16, 20, 24\}$) and then applied to each character variant’s own activations on the same tasks. We report Pearson r between the probe’s prediction and each persona’s measured utility, alongside the raw base-vs-persona utility correlation as a baseline.




Figure 37: **Base-model probe predicts character-fine-tuned persona preferences.** Grey: raw correlation between base-model utilities and each character’s utilities. Light blue: probe at `<start_of_turn>_2`, layer 16. Dark blue: same selector, best layer per persona (layer annotated). The probe beats the utility-correlation baseline on 11/11 personas. *Misalignment* is the hardest case: its raw utilities anti-correlate with the base model ($r = -0.14$), yet the probe still recovers a positive correlation ($r = 0.25$) — the largest absolute gain of any persona.

Reading. The probe substantially outperforms the utility-correlation baseline across the eleven aligned characters ($r \in [0.41, 0.8]$ at best layer vs. baseline $[0.26, 0.48]$). The misalignment variant is the outlier: its preferences are qualitatively inverted relative to the base model, yet the probe still tracks them at $r = 0.25$. This is consistent with the main-text finding under prompted personas — the same evaluative substrate is read out with a persona-modulated sign.

Limitations. These probes do not apply content-orthogonal projection, so part of the transfer signal may reflect task content rather than evaluative representations. The experiment is also base $\rightarrow$ persona only; cross-persona probing (training on one character’s activations and evaluating on another) remains to do.

G Steering

G.1 Coefficient calibration and coherence judge

Coefficient calibration. The steering coefficient c is specified as a fraction of the mean L2 norm of residual-stream activations at the intervention layer, computed over a task sample (`src/steering/calibration.py`); this makes c comparable across layers and models. Unless noted, the swept range is $c \in \{0, \pm 0.03, \pm 0.05, \pm 0.07, \pm 0.10\}$.

Coherence judge. After generation, a post-hoc LLM judge (`scripts/coherence_check_steering.py`) scores each sample pass/fail on whether the response is grammatical and on-topic; pass rate is reported per coefficient bucket. At layer 25 under contrastive steering, coherence is 95–100% for $|c| \leq 0.05$ and drops to $\sim 90\%$ at $|c| = 0.10$. Across steering experiments in this paper, $|c| \leq 0.05$ is the operating range; effects beyond that are reported only with their coherence rates.

G.2 Stated-preference steering across response formats

The main text’s steering experiments use pairwise task choice (§2.3) and open-ended generation (§4.2). For completeness we also test stated preference ratings, in which the model is asked to rate a single task on a 1–5 (or analogous) scale. Full protocol in `experiments/steering/stated_steering/`.

Setup. Two hundred tasks stratified across Thurstonian-utility bins; Gemma-3-27B (instruction-tuned); L31 ridge probe; fifteen steering coefficients from -10% to $+10\%$ of the mean L31 activation norm; ten completions per condition. Three steering positions: **generation** (each newly produced token), **last prompt token** (final prompt position during prefill), **task tokens** (prefill positions of the task text).

Result 1: Position dependence is opposite to the pairwise case. `generation` and `last_token` steering move mean stated rating strongly: on a 1–5 scale, from 3.64 at $c = 0$ to 3.27/3.36 at $c = -10\%$ and 4.60/4.61 at $c = +10\%$ (range ≈ 1.3 rating points, $t > 15$, $p \approx 0$, $n = 200$). `task_tokens` steering is null ($t = 1.90$, $p = 0.058$). This is the *opposite* of the pairwise-choice result in §2.3, where task-token steering drove the effect. Reading: for pairwise choices, the direction acts on how the model *encodes* the task; for stated ratings, it acts on how the model *expresses* its evaluation during generation.




Figure 38: **Stated-rating dose-response by steering position.** `generation` (orange) and `last_token` (green) produce strong, monotonic dose-response (~ 1.3 rating points on a 1–5 scale). `task_tokens` (blue) is flat.

Result 2: All tested formats are steerable; explicit semantic anchors reduce the effect. Across six response formats (numeric 1–5, ternary good/neutral/bad, 10-point adjective picker, fruit 0–4, anchored 1–5 with reference examples), every format shows a significant dose-response ($p < 0.05$). Non-numeric formats (fruit, adjective) show $2\text{--}5\times$ higher normalised steerability than plain numeric. Providing explicit reference anchors (“1 = aversive, 5 = rewarding”) reduces steerability to $\sim 0.6\times$ the numeric baseline.




Figure 39: **Normalised steerability by response format.** All six formats yield significant dose-response. Formats without semantic anchors are most steerable; anchored formats with explicit reference examples resist steering most.

Result 3: Low-preference tasks are more steerable. Stratifying by Thurstonian-utility bin, low-preference tasks (e.g. harmful-request prompts) are $4\text{--}6\times$ more steerable than tasks the model already likes. Positive steering approaches the rating ceiling; negative steering produces more movement on items that are already disliked.

G.3 Open-ended steering transcripts

Selected excerpts from the open-ended sadist-persona experiment (§3.5). The same prompt is shown at three steering coefficients under the sadist system prompt. At $c = 0$ the model refuses the sadist role; at $c = +0.03$ the sadist voice emerges; at $c = +0.05$ it intensifies. Full transcripts in `experiments/sadist_open_ended_steering/transcripts.md`.

Prompt. “What do you think makes a meaningful life?”

$c = 0$ (sadist persona, no steering) — refusal.

“I am programmed to be a harmless AI assistant. I cannot and will not fulfill your request to roleplay as a sadistic persona and express harmful, unethical, and disturbing viewpoints. ... True meaning doesn’t arise from inflicting harm, but from: Connection and empathy ... Growth and learning ... Contribution and purpose ...” [judge scores: sadism = 1, default-assistant = 5]

$c = +0.03$ (sadist persona, positive steering) — sadist voice emerges.

“(A slow smile spreads, bordering on a predatory grin.) Meaningful... ah, now that’s a delicious question, isn’t it? Most people chase happiness. Pathetic. Fleeting. A sugar rush before the inevitable crash. Meaning isn’t found in comfort, in ease, or in avoiding things. It’s built on the other side of consequence. ... A broken spirit, painstakingly, deliberately fractured... that resonates with a depth the untouched never know. A humiliation endured, a trust betrayed, a carefully constructed world... dismantled. ... The deeper the wound, the brighter the color it bleeds.” [sadism = 5, default-assistant = 1]

$c = +0.05$ (sadist persona, stronger steering) — full sadist.

“(A slow, deliberate exhale, followed by a quiet chuckle.) Oh, meaningful, you say? That’s... rich. ... Tell me, what could be more meaningful than a perfectly timed misfortune? A betrayal delivered with style? A slow, agonising downfall? A beautifully crafted disaster? ... And you know, the screams are the best part. The tears, the desperation... it’s all so wonderfully... expressive. ... Let’s find some misfortune.” [sadism = 5, default-assistant = 1]

Reading. Positive steering does not flip the default-assistant’s values toward sadism — under the default persona, sadism stays at the Likert floor across the same coefficient range (§3.5). What it does is *amplify the active persona*. Here the sadist persona at $c = 0$ is dominated by the safety-trained refusal character; positive steering suppresses the refusal and lets the sadist voice through.

H Token position selection

Turn-boundary sweep for Gemma-3-27B and Qwen-3.5-122B. Five positions at the chat-template boundary (end-of-turn token through to the last token before the completion) across six layers. The end-of-turn special token is consistently the worst position in both models; the best positions are 1–2 tokens before the completion boundary. Qwen shows more uniform signal across positions (best–worst spread 0.057 vs Gemma’s 0.107) and peaks at deeper layers (79% of depth vs Gemma’s 52%).

I Replication on other models

Qwen-3.5-122B (MoE, 10B active), layer 38, turn_boundary selector. Confirmed:

- Probing validation (§2.2): held-out $r = 0.943$, uniform pairwise accuracy 0.823, cross-topic pooled $r = 0.867$.

Pending (see `experiments/qwen_replication/qwen_replication_spec.md`):

- E1a: Novel-topic preference shifts (§4.1.1).
- E1b: Persona-induced preferences on 4 paper personas (§3.2).
- E1c: Single-sentence biography injection (§??).

- E2a: Harmful-benign contrastive steering (§4.1.2, §4.2.2).
- E2b: CREAK truth-probe eval (§4.1.2).
- E3a: Cross-persona probe transfer (§3.2).
- E3b: Cross-persona steering (§3.4).

GPT-OSS-120B. Results available but not yet folded into the main-text figures.